

# Cloud and Machine Learning

## Project 1: "ImageNet on BareMetal/Cloud"

### **PART 1 : Running ImageNet on Bare metal**

The steps followed have been listed below:

1. Access the Prince Cluster

Use an SSH Client to connect to the Prince Cluster:

- We used PuTTY to connect to host gw.hpc.nyu.edu on port 22 using SSH protocol.
- Enter your NetID password and you should be authorized on the gateway server.
- Use the following ssh command to connect to the prince cluster.  
ssh prince.hpc.nyu.edu


2. Get access to the ImageNet dataset that we will be using for the project. The dataset can be found on the prince cluster /scratch/work/public/imagenet
3. We proceed to create a small subset of the training data for our use under /home/net-id  
We divide the data into train and val subsets that will be used for the training and validation of the model we create.
4. Get the Imagenet example from pytorch.  
git clone <https://github.com/pytorch/examples>
5. Move to the imagenet folder containing the python code.  
cd examples → cd imagenet
6. Get an interactive node  
Login into frontend node. ssh prince1.hpc.nyu.edu  
srun -t10:00:00 --gres=gpu:1 --mem 102400 --pty /bin/bash
7. Load module. We are using the stable version.  
module load python3/intel/3.6.3 cuda/9.0.176 nccl/cuda9.0/2.4.2
8. Setup the python virtual environment.
  - mkdir pytorch\_env
  - cd pytorch\_env
  - virtualenv --system-site-packages py3.6.3
  - source py3.6.3/bin/activate

- Install pytorch  
pip3 install [http://download.pytorch.org/whl/cu92/torch-0.4.1-cp36-cp36m-linux\\_x86\\_64.whl](http://download.pytorch.org/whl/cu92/torch-0.4.1-cp36-cp36m-linux_x86_64.whl)
  - Install torchvision  
pip3 install torchvision
9. Activate the virtual environment.  
source ~/pytorch\_env/py3.6.3/bin/activate
  10. Run the job. To train the model we run 'main.py' with the desired model architecture(Alexnet) and a path to the Imagenet dataset.  
python /home/as13594/examples/imagenet/main.py -a alexnet -b 8 --epochs 1 --lr 0.01 /home/as13594/
  11. Use nvprof for profiling the performance and run the job on the reserved GPU nodes using the following command  
srun --reservation=chung --gres=gpu:1 --time=01:00:00 --gres=gpu:p40:4 --cpus-per-task=28  
nvprof python /home/as13594/examples/imagenet/main.py -a alexnet -b 20 --epochs 1 --lr 0.01 /home/as13594/

## PART 2 : Running ImageNet on Cloud

For the second part of the project we run the code on a Google Colab notebook on the AWS Cloud Computing Platform.

1. We begin with login onto the Google Colaboratory.
2. Upload the python file containing the source code that needs to be run on the drive.
3. Create a notebook instance.
4. Once the notebook instance has been created, we mount the drive containing the data(testing and validation) to be used for building the model.  
from google.colab import drive  
drive.mount('/content/gdrive')
5. Change runtime type to GPU.
6. Run the code.  
python /home/as13594/examples/imagenet/main.py -a alexnet -b 8 --epochs 1 --lr 0.01 /home/as13594/

Active sessions			
Title		Last execution	RAM used
 CloudAndML.ipynb Current session	GPU	0 minutes ago	0.91 GB
			<a href="#">TERMINATE</a>

## Implementation and Analysis:

As a part of this project we attempt to train a model with the desired model architecture(Alexnet) on the Imagenet dataset using pytorch and run it on both BareMetal and Cloud platforms to access its relative performance on both platforms.

## 1. Usability:

- **BareMetal**- A bare metal server is a single tenant physical server. Since they offer single-tenant environments i.e. a single servers' physical resources may not be shared between two or more tenants, they can be used to run dedicated services without any interruptions for longer durations. Bare-metal servers offer isolation and are free of the “noisy neighbor” effect that plagues virtual environments. Network latency is minimized for better performance, and the tenant enjoys root access. Bare metal is highly customizable, and the tenant may optimize the server based upon their individual needs.  
Bare metal servers do not require the use of several layers of software, unlike the virtual environment, which has at least one additional layer of software – a Type 1 hypervisor. This implies that there is one less layer of software between the user and the physical hardware in everyday use. Hence, we can expect better performance.
- **Cloud** : Cloud offers a distributed environment comprised of multi-tenant, virtualized servers. The host machine shares its resources with multiple virtual instances. They each get a portion of CPU, RAM and storage. Cloud Servers allow you to add resources to individual virtual machines (vertical scaling) or add whole new servers (horizontal scaling) at any time, in a matter of minutes. This scalability makes Cloud Servers better suited to variable workloads, where the ability to dynamically scale performance is more important than sheer horsepower. Cloud computing environments are more prone to latency for various reasons. For example, if VMs are on separate networks, it can lead to packet delays. With cloud environments, you do not have a direct connection with the physical hardware, as there is a hypervisor layer between your app and physical resources. Thus, the chances are that VMs will suffer from a higher latency than if you were running apps directly on a bare metal server. Furthermore, performance bottlenecks may occur due to the sheer number of tenants. If you have noisy neighbors who like to run resource-intensive workloads on their share of the server, they may very well impact you leading to degraded performance.

### Analysis of Usability between BareMetal and Cloud:

Running our program on both BareMetal and Cloud environments gave us a chance to explore the pros and cons of both domains.

- **Google Collaboratory offered a more friendly user interface. It was relatively easy to mount data and process it on a GPU using a notebook. There was no need for the user to manage the dependencies/configuration of the execution environment.**
- **In case of BareMetal, we had to set up our own virtual execution environment before we were able to run any jobs on the GPU. It was relatively more tedious to transfer the data to the server and to run the jobs.**
- **GPU based cloud computing provides the ease of scalability as opposed to BareMetal Servers. This makes them fit for variable workloads.**
- **Since resources in the Bare Metal environments are dedicated to users this could lead to under-utilization of resources as opposed to cloud environments that offer better utilization of resources.**
- **Bare-metal environments are single tenant systems and are comparatively more secure.**


## 2. Hyperparameter settings:

- **Learning Rate** : The weights of a neural network cannot be calculated using an analytical method. Instead, the weights must be discovered via an empirical optimization procedure called gradient descent.

The learning rate is a hyperparameter that controls how much to change the model in response to the estimated error each time the model weights are updated. This parameter scales the magnitude of our weight updates in order to minimize the network's loss function. Choosing the learning rate is challenging as a value too small may result in a long training process that could get stuck, whereas a value too large may result in learning a sub-optimal set of weights and too fast or unstable training process.

We might start with a large value like 0.1, then try exponentially lower values: 0.01, 0.001, etc. The training should start from a relatively large learning rate because, in the beginning, random weights are far from optimal, and then the learning rate can decrease during training to allow more fine-grained weight updates. Training with a smaller learning rate would allow more fine-grained weight updates.

Note: We use 0.01 as the initial learning rate for AlexNet and then change it to .05 and .0001

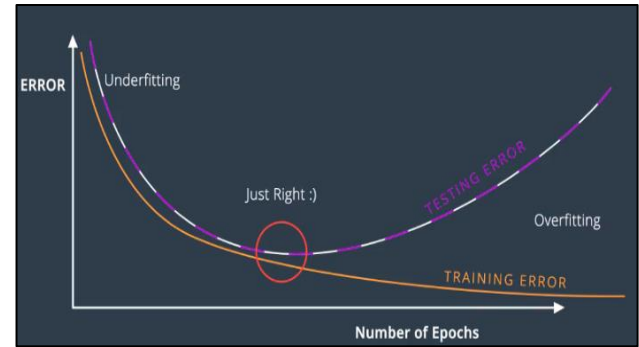


- **BATCH SIZE:** Batch size is used in machine learning to refer to the number of training examples utilized in one iteration. Choosing a batch size that is too small will introduce a high degree of variance within each batch as it is unlikely that a small sample is a good representation of the entire dataset. Generally, the larger the batch size the faster the model will complete each epoch in training. However, the tradeoff is that the quality of the model may degrade as we increase the batch size. If a batch size is too large, it may not fit in memory of the compute instance used for training and it will have the tendency to overfit the data.

Lowering the learning rate and decreasing the batch size will allow the network to train better, especially in the case of fine-tuning.

- **EPOCH:** The number of epochs is a hyperparameter of gradient descent that controls the number of complete passes through the training dataset. The number of epoch will decide- how many times we will change the weights of the network.

As the number of epochs increases, the number of times weights are changed in the neural network increases and the boundary goes from underfitting to optimal to overfitting.



### 3. Performance:

#### BareMetal vs. Cloud Platforms: At a Glance

Deep learning consumes huge amount of computing capacity, because it often relies on deeper networks and larger training datasets to improve the model accuracy

Our experiments focuses on following performance metrics:

- (1) **Throughput** measured by the number of images processed per second during the training steps, and
- (2) **Time elapsed:** measured by the total execution time from the time a training job is launched until it finished.

In other words, the time elapsed metric includes the performance impact of data loading and pre-processing while the throughput metric does not thereby justifying our choice of metrics.

#### Performance Analysis:

##### Example Run 1

- **Running on Cloud:**

We begin with training the model with the following hyperparameter settings: Batch size - 20 Epoch -1 and the Learning Rate has been set to 0.01. It takes approximately 221 iterations and processes images in batches of 20 and passes over the dataset one time.

Training the model with the following hyperparameter settings: Batch size - 20 Epoch -1 and Learning Rate set to 0.01

```
!time nvprof python main.py -a alexnet -b 20 --epochs 1 --lr 0.01 Imagenet
```

```
==597== NVPROF is profiling process 597, command: python3 main.py -a alexnet -b 20 --epochs 1 --lr 0.01 Imagenet
=> creating model 'alexnet'
Epoch: [0][ 0/221] Time 1.336 ( 1.336) Data 1.015 ( 1.015) Loss 6.9105e+00 (6.9105e+00) Acc@1 0.00 ( 0.00) Acc@5 0.00 ( 0.00)
Epoch: [0][ 10/221] Time 0.041 ( 0.198) Data 0.001 ( 0.143) Loss 6.7353e+00 (6.8455e+00) Acc@1 35.00 ( 19.55) Acc@5 100.00 ( 65.00)
Epoch: [0][ 20/221] Time 0.421 ( 0.178) Data 0.407 ( 0.136) Loss 4.5311e+00 (6.6527e+00) Acc@1 30.00 ( 22.62) Acc@5 80.00 ( 80.71)
Epoch: [0][ 30/221] Time 0.460 ( 0.173) Data 0.446 ( 0.136) Loss 4.4275e+00 (6.3101e+00) Acc@1 25.00 ( 22.58) Acc@5 85.00 ( 79.35)
Epoch: [0][ 40/221] Time 0.029 ( 0.159) Data 0.000 ( 0.125) Loss 4.4720e+00 (5.8564e+00) Acc@1 20.00 ( 22.68) Acc@5 70.00 ( 79.39)
Epoch: [0][ 50/221] Time 0.641 ( 0.166) Data 0.616 ( 0.133) Loss 2.6310e+00 (5.5344e+00) Acc@1 25.00 ( 22.75) Acc@5 100.00 ( 80.59)
Epoch: [0][ 60/221] Time 0.043 ( 0.155) Data 0.001 ( 0.122) Loss 1.5077e+00 (4.9519e+00) Acc@1 30.00 ( 24.26) Acc@5 100.00 ( 83.77)
Epoch: [0][ 70/221] Time 0.220 ( 0.154) Data 0.205 ( 0.123) Loss 1.3294e+00 (4.4341e+00) Acc@1 45.00 ( 26.55) Acc@5 100.00 ( 86.06)
Epoch: [0][ 80/221] Time 0.471 ( 0.158) Data 0.455 ( 0.127) Loss 1.1508e+00 (4.0449e+00) Acc@1 25.00 ( 27.47) Acc@5 100.00 ( 87.78)
Epoch: [0][ 90/221] Time 0.040 ( 0.154) Data 0.000 ( 0.123) Loss 1.3185e+00 (3.7233e+00) Acc@1 35.00 ( 29.07) Acc@5 100.00 ( 89.12)
Epoch: [0][100/221] Time 0.400 ( 0.155) Data 0.385 ( 0.125) Loss 9.9271e-01 (3.4642e+00) Acc@1 45.00 ( 30.69) Acc@5 100.00 ( 90.20)
Epoch: [0][110/221] Time 0.168 ( 0.153) Data 0.154 ( 0.123) Loss 1.1549e+00 (3.2504e+00) Acc@1 25.00 ( 31.62) Acc@5 100.00 ( 91.08)
Epoch: [0][120/221] Time 0.211 ( 0.151) Data 0.192 ( 0.122) Loss 1.0894e+00 (3.0752e+00) Acc@1 60.00 ( 33.14) Acc@5 100.00 ( 91.82)
Epoch: [0][130/221] Time 0.149 ( 0.150) Data 0.132 ( 0.121) Loss 9.1729e-01 (2.9230e+00) Acc@1 65.00 ( 34.01) Acc@5 100.00 ( 92.44)
Epoch: [0][140/221] Time 0.168 ( 0.149) Data 0.150 ( 0.120) Loss 1.2687e+00 (2.7844e+00) Acc@1 45.00 ( 35.50) Acc@5 100.00 ( 92.98)
Epoch: [0][150/221] Time 0.314 ( 0.148) Data 0.296 ( 0.119) Loss 1.1607e+00 (2.6707e+00) Acc@1 50.00 ( 36.36) Acc@5 100.00 ( 93.44)
Epoch: [0][160/221] Time 0.035 ( 0.147) Data 0.002 ( 0.118) Loss 1.0375e+00 (2.5662e+00) Acc@1 40.00 ( 37.42) Acc@5 100.00 ( 93.85)
Epoch: [0][170/221] Time 0.414 ( 0.148) Data 0.398 ( 0.119) Loss 9.1133e-01 (2.4695e+00) Acc@1 40.00 ( 38.42) Acc@5 100.00 ( 94.21)
Epoch: [0][180/221] Time 0.067 ( 0.146) Data 0.046 ( 0.117) Loss 7.7018e-01 (2.3863e+00) Acc@1 70.00 ( 39.64) Acc@5 100.00 ( 94.53)
Epoch: [0][190/221] Time 0.073 ( 0.146) Data 0.155 ( 0.117) Loss 1.0964e+00 (2.3116e+00) Acc@1 50.00 ( 40.65) Acc@5 100.00 ( 94.82)
Epoch: [0][200/221] Time 0.343 ( 0.146) Data 0.323 ( 0.117) Loss 1.2081e+00 (2.2559e+00) Acc@1 45.00 ( 40.77) Acc@5 100.00 ( 95.07)
Epoch: [0][210/221] Time 0.032 ( 0.148) Data 0.003 ( 0.120) Loss 1.0584e+00 (2.2049e+00) Acc@1 55.00 ( 40.95) Acc@5 100.00 ( 95.31)
Epoch: [0][220/221] Time 0.139 ( 0.145) Data 0.000 ( 0.116) Loss 9.2850e-01 (2.1576e+00) Acc@1 60.00 ( 41.38) Acc@5 100.00 ( 95.51)
Test: [ 0/41] Time 1.142 ( 1.142) Loss 4.9381e-01 (4.9381e-01) Acc@1 95.00 ( 95.00) Acc@5 100.00 (100.00)
Test: [10/41] Time 0.012 ( 0.252) Loss 4.6495e-01 (3.9548e-01) Acc@1 95.00 ( 94.55) Acc@5 100.00 (100.00)
Test: [20/41] Time 0.231 ( 0.201) Loss 7.2971e-01 (4.3935e-01) Acc@1 60.00 ( 93.10) Acc@5 100.00 (100.00)
Test: [30/41] Time 0.021 ( 0.188) Loss 1.3648e+00 (6.9853e-01) Acc@1 15.00 ( 76.45) Acc@5 100.00 (100.00)
Test: [40/41] Time 0.092 ( 0.168) Loss 1.8127e+00 (8.9963e-01) Acc@1 10.00 ( 59.63) Acc@5 100.00 (100.00)
* Acc@1 59.630 Acc@5 100.000
==597== Profiling application: python3 main.py -a alexnet -b 20 --epochs 1 --lr 0.01 Imagenet
==597== Profiling result:
Type Time(%) Time Calls Avg Min Max Name
GPU activities: 31.82% 2.61611s 14632 178.79us 1.0240us 1.9348ms _ZN2at6native29vectorized_elementwise_kernelIi4EZZNS0_15add_kernel_cudaE
10.47% 861.15ms 2219 388.08us 76.958us 1.3111ms volta_sgemv_128x64_nt
7.64% 627.96ms 1280 490.60us 199.48us 1.3422ms volta_sgemv_128x64_nn
5.50% 452.33ms 4568 99.022us 1.0230us 1.3371ms _ZN2at6native29vectorized_elementwise_kernelIi4EZZNS0_23gpu_kernel_with_sc
4.63% 380.84ms 786 484.53us 107.04us 1.2109ms volta_sgemv_128x32_slided1x4_tn
3.85% 316.17ms 544 581.19us 1.0870us 32.267ms [CUDA memcopy HtoD]
```

## Performance Profiling using nvprof:

Profiling the model with the following hyperparameter settings: Batch size - 20 Epoch -1 and Learning Rate set to 0.01

```
!time nvprof python main.py -a alexnet -b 20 --epochs 1 --lr 0.01 Imagenet > log1.txt
```

```
==371== NVPROF is profiling process 371, command: python3 main.py -a alexnet -b 20 --epochs 1 --lr 0.01 Imagenet
==371== Profiling application: python3 main.py -a alexnet -b 20 --epochs 1 --lr 0.01 Imagenet
==371== Profiling result:
Type Time(%) Time Calls Avg Min Max Name
GPU activities: 24.21% 2.57984s 14632 176.32us 1.3120us 1.9925ms _ZN2at6native29vectorized_elementwise_kernelIi4EZZNS0_15add_kernel_cudaERNIS_14TensorIteratorEN3c
15.71% 1.67423s 2219 754.50us 84.735us 1.3128ms volta_sgemv_128x64_nt
10.92% 1.16393s 1280 909.32us 251.36us 1.3431ms volta_sgemv_128x64_nn
4.80% 510.96ms 786 650.08us 113.76us 1.2119ms volta_sgemv_128x32_slided1x4_tn
4.18% 445.15ms 4568 97.450us 1.3120us 1.3940ms _ZN2at6native29vectorized_elementwise_kernelIi4EZZNS0_23gpu_kernel_with_scalarsIZZNS0_15mul_kernel
3.56% 379.66ms 660 575.25us 84.799us 1.1280ms volta_sgemv_128x32_slided1x4_nn
3.45% 368.04ms 224 1.6430ms 799.86us 1.7007ms volta_scudnn_128x64_stridedB_splitK_medium_nn_v1
3.02% 321.83ms 544 591.60us 1.3760us 32.270ms [CUDA memcopy HtoD]
2.81% 298.89ms 265 1.1279ms 277.88us 1.1548ms volta_scudnn_128x64_relu_xrgs_large_nn_v1
2.61% 278.56ms 663 420.16us 45.087us 673.98us void at::native::GLOBAL__N_63_tmpxft_00002580_00000000_10_DilatedMaxPool2d_compute_75_cpp1_iidb9
2.49% 265.24ms 4625 57.349us 1.0560us 653.14us _ZN2at6native29vectorized_elementwise_kernelIi4EZZNS0_16fill_kernel_cudaERNIS_14TensorIteratorEN3c
1.89% 201.84ms 1458 136.44us 40.416us 207.77us void cudnn::wInograd_nonfused::wInogradForwardData4x4x4_float, float>(cudnn::wInograd_nonfused::wInog
1.74% 185.05ms 3381 54.73us 2.7840us 189.15us _ZN2at6native29vectorized_elementwise_kernelIi4EZZNS0_21threshold_kernel_implfEEERNIS_14TensorItera
1.71% 182.15ms 1458 124.93us 43.999us 172.80us void cudnn::wInograd_nonfused::wInogradForwardOutput4x4x4_float, float>(cudnn::wInograd_nonfused::wIn
1.65% 175.28ms 411 426.46us 1.6320us 62.590ms [CUDA memcopy DtoH]
1.63% 173.36ms 711 243.82us 66.527us 447.45us void cudnn::wInograd_nonfused::wInogradForwardData9x9_5x5_float, float>(cudnn::wInograd_nonfused::w
1.23% 130.57ms 669 195.18us 87.999us 239.20us void cudnn::wInograd_nonfused::wInogradGradOutput4x4x4_float, float>(cudnn::wInograd_nonfused::wInog
1.13% 119.87ms 488 245.64us 60.447us 352.67us void cudnn::wInograd_nonfused::wInogradForwardOutput9x9_5x5_float, float>(cudnn::wInograd_nonfused:
1.06% 112.73ms 1310 86.049us 18.112us 164.77us _ZN2at6native27unrolled_elementwise_kernelIiZZNS0_15add_kernel_cudaERNIS_14TensorIteratorEN3c106scal
1.06% 112.42ms 1458 77.103us 58.528us 116.06us void cudnn::wInograd_nonfused::wInogradForwardFilter4x4x4_float, float>(cudnn::wInograd_nonfused::wIn
0.76% 80.909ms 786 102.94us 10.975us 156.57us void at::native::GLOBAL__N_63_tmpxft_00002580_00000000_10_DilatedMaxPool2d_compute_75_cpp1_iidb9
0.75% 79.923ms 223 358.40us 145.82us 374.91us void cudnn::wInograd_nonfused::wInogradGradDelta9x9_5x5_float, float>(cudnn::wInograd_nonfused::wIn
0.72% 76.321ms 1629 46.851us 3.1040us 91.039us _ZN2at6native13reduce_kernelIi15I2EliENS0_8ReduceOpIfNIS0_14func_wrapper_tIFZNIS0_15sum_kernel_implI
0.59% 62.924ms 20 3.1462ms 506.46us 5.2813ms volta_ggemv_32x32_nt
0.59% 62.511ms 669 93.439us 18.912us 139.26us void cudnn::wInograd_nonfused::wInogradGradDelta4x4x4_float, float>(cudnn::wInograd_nonfused::wInog
0.54% 57.681ms 669 86.220us 12.864us 130.65us void cudnn::wInograd_nonfused::wInogradGradData4x4x4_float, float>(cudnn::wInograd_nonfused::wInog
0.43% 45.828ms 262 174.92us 45.504us 182.37us void at::native::GLOBAL__N_69_tmpxft_00001db3_00000000_10_AdaptiveAveragePooling_compute_75_cpp1_
0.34% 36.598ms 221 165.60us 41.311us 188.13us void at::native::GLOBAL__N_69_tmpxft_00001db3_00000000_10_AdaptiveAveragePooling_compute_75_cpp1_
0.34% 35.969ms 16 2.2481ms 682.46us 4.7935ms void cudnn::detail::dgrad_engine_float, int=128, int=6, int=8, int=3, int=3, int=5, bool=1>(int, in
0.32% 33.996ms 12 2.8330ms 1.5310ms 5.0412ms void cudnn::detail::implicit_convolve_sgemv_float, float, int=512, int=6, int=8, int=3, int=3, int=
0.31% 33.057ms 488 67.739us 62.559us 73.055us void cudnn::wInograd_nonfused::wInogradForwardFilter9x9_5x5_float, float>(cudnn::wInograd_nonfused:
0.25% 26.578ms 12 2.2148ms 574.07us 5.3919ms void cudnn::detail::wInograd_algo_engine_float, int=128, int=6, int=8, int=3, int=3, int=5, bool=1, in
0.24% 25.668ms 28 916.71us 324.28us 1.4948ms volta_cgmmv_32x32_tn
0.23% 24.202ms 9 2.6891ms 2.0869ms 3.4440ms void fft2d_r2c_32x32_float, bool=0, unsigned int=1, bool=1>(float2*, float const *, int, int, int,
```

```
real 5m17.019s
user 0m58.431s
sys 0m13.777s
```

The time taken to finish training for this run is very less as could be expected from because we started with a greater value of learning rate which determines the step size. One can also note that the accuracy is just 25%.

- **Running on BareMetal:**

We train the model with the same hyperparameter settings: Batch size - 20 Epoch -1 and the Learning Rate has been set to 0.01. It takes approximately 221 iterations and processes images in batches of 20 and passes over the dataset one time.

```
(py3.6.3) [asl3594@gpu-61 ~]$ time nvprof python /home/asl3594/examples/imagenet/main.py -a alexnet -b 20 --epochs 1 --lr 0.01 /home/asl3594/
==48297== NVPROF is profiling process 48297, command: python /home/asl3594/examples/imagenet/main.py -a alexnet -b 20 --epochs 1 --lr 0.01 /home/asl3594/
=> creating model 'alexnet'
Epoch: [0] [ 0/221] Time 1.109 ( 1.109) Data 0.886 ( 0.886) Loss 6.9065e+00 (6.9065e+00) Acc@1 0.00 ( 0.00) Acc@5 0.00 ( 0.00)
Epoch: [0] [10/221] Time 0.039 ( 0.209) Data 0.000 ( 0.167) Loss 6.6925e+00 (6.8328e+00) Acc@1 20.00 ( 14.09) Acc@5 100.00 ( 65.91)
Epoch: [0] [20/221] Time 0.465 ( 0.183) Data 0.454 ( 0.152) Loss 6.3945e+00 (6.6280e+00) Acc@1 25.00 ( 16.19) Acc@5 75.00 ( 77.62)
Epoch: [0] [30/221] Time 0.026 ( 0.174) Data 0.000 ( 0.143) Loss 6.2642e+00 (7.5265e+00) Acc@1 25.00 ( 19.52) Acc@5 70.00 ( 76.94)
Epoch: [0] [40/221] Time 0.046 ( 0.161) Data 0.010 ( 0.130) Loss 1.1596e+01 (7.2846e+00) Acc@1 20.00 ( 20.73) Acc@5 25.00 ( 74.02)
Epoch: [0] [50/221] Time 0.028 ( 0.166) Data 0.000 ( 0.135) Loss 5.0911e+00 (6.9655e+00) Acc@1 35.00 ( 20.88) Acc@5 75.00 ( 74.22)
Epoch: [0] [60/221] Time 0.046 ( 0.159) Data 0.002 ( 0.128) Loss 4.8199e+00 (6.6354e+00) Acc@1 25.00 ( 21.39) Acc@5 70.00 ( 74.43)
Epoch: [0] [70/221] Time 0.039 ( 0.161) Data 0.000 ( 0.131) Loss 3.6673e+00 (6.2894e+00) Acc@1 25.00 ( 22.39) Acc@5 80.00 ( 74.58)
Epoch: [0] [80/221] Time 0.047 ( 0.155) Data 0.003 ( 0.125) Loss 1.8776e+00 (5.9031e+00) Acc@1 20.00 ( 22.47) Acc@5 100.00 ( 76.98)
Epoch: [0] [90/221] Time 0.046 ( 0.156) Data 0.000 ( 0.125) Loss 1.4587e+00 (5.4372e+00) Acc@1 35.00 ( 22.86) Acc@5 100.00 ( 79.51)
Epoch: [0] [100/221] Time 0.042 ( 0.151) Data 0.015 ( 0.121) Loss 1.3594e+00 (5.0451e+00) Acc@1 25.00 ( 23.12) Acc@5 100.00 ( 81.53)
Epoch: [0] [110/221] Time 0.034 ( 0.153) Data 0.000 ( 0.123) Loss 1.4235e+00 (4.7189e+00) Acc@1 30.00 ( 23.29) Acc@5 100.00 ( 83.20)
Epoch: [0] [120/221] Time 0.050 ( 0.152) Data 0.005 ( 0.122) Loss 1.4973e+00 (4.4490e+00) Acc@1 25.00 ( 23.72) Acc@5 100.00 ( 84.59)
Epoch: [0] [130/221] Time 0.035 ( 0.152) Data 0.005 ( 0.122) Loss 1.4152e+00 (4.2217e+00) Acc@1 20.00 ( 23.97) Acc@5 100.00 ( 85.76)
Epoch: [0] [140/221] Time 0.036 ( 0.150) Data 0.002 ( 0.120) Loss 1.4597e+00 (4.0232e+00) Acc@1 15.00 ( 24.04) Acc@5 100.00 ( 86.77)
Epoch: [0] [150/221] Time 0.043 ( 0.150) Data 0.000 ( 0.120) Loss 1.4626e+00 (3.8516e+00) Acc@1 15.00 ( 23.77) Acc@5 100.00 ( 87.65)
Epoch: [0] [160/221] Time 0.035 ( 0.150) Data 0.000 ( 0.120) Loss 1.4795e+00 (3.7016e+00) Acc@1 20.00 ( 23.66) Acc@5 100.00 ( 88.42)
Epoch: [0] [170/221] Time 0.049 ( 0.150) Data 0.000 ( 0.120) Loss 1.4996e+00 (3.5724e+00) Acc@1 15.00 ( 23.80) Acc@5 100.00 ( 89.09)
Epoch: [0] [180/221] Time 0.050 ( 0.149) Data 0.002 ( 0.119) Loss 1.4171e+00 (3.4577e+00) Acc@1 30.00 ( 24.14) Acc@5 100.00 ( 89.70)
Epoch: [0] [190/221] Time 0.027 ( 0.150) Data 0.000 ( 0.120) Loss 1.5100e+00 (3.3522e+00) Acc@1 20.00 ( 24.40) Acc@5 100.00 ( 90.24)
Epoch: [0] [200/221] Time 0.050 ( 0.148) Data 0.008 ( 0.118) Loss 1.4543e+00 (3.2566e+00) Acc@1 30.00 ( 24.23) Acc@5 100.00 ( 90.72)
Epoch: [0] [210/221] Time 0.047 ( 0.148) Data 0.000 ( 0.118) Loss 1.3781e+00 (3.1688e+00) Acc@1 30.00 ( 24.19) Acc@5 100.00 ( 91.16)
Epoch: [0] [220/221] Time 0.096 ( 0.146) Data 0.000 ( 0.116) Loss 1.4402e+00 (3.0951e+00) Acc@1 40.00 ( 24.18) Acc@5 100.00 ( 91.53)
Test: [ 0/40] Time 1.187 ( 1.187) Loss 1.3724e+00 (1.3724e+00) Acc@1 20.00 ( 20.00) Acc@5 100.00 (100.00)
Test: [10/40] Time 0.009 ( 0.238) Loss 1.1259e+00 (1.3368e+00) Acc@1 100.00 ( 25.91) Acc@5 100.00 (100.00)
Test: [20/40] Time 0.246 ( 0.200) Loss 1.6382e+00 (1.2612e+00) Acc@1 0.00 ( 56.43) Acc@5 100.00 (100.00)
Test: [30/40] Time 0.006 ( 0.190) Loss 1.4159e+00 (1.3745e+00) Acc@1 0.00 ( 38.23) Acc@5 100.00 (100.00)
* Acc@1 29.625 Acc@5 100.000
==48297== Profiling application: python /home/asl3594/examples/imagenet/main.py -a alexnet -b 20 --epochs 1 --lr 0.01 /home/asl3594/
==48297== Profiling result:
Type Time(%) Time Calls Avg Min Max Name
GPU activities: 23.74% 1.17533s 14630 80.337us 1.3110us 831.77us _ZN2at6native29vectorized_elementwise_kernelILi4EZZNS0_15add_kernel_cudaERNNS
lvE_cIEvEUKulvE2_cIEvEULffe_NS_6detail5ArrayIPcLi3EEEEvTi0_T1
8.52% 421.93ms 1338 315.35us 64.543us 529.15us maxwell_sgemv_128x64_nt
7.22% 357.60ms 1012 353.36us 122.37us 525.50us maxwell_sgemv_128x64_nn
6.76% 334.74ms 663 504.88us 52.320us 805.44us void at::native::GLOBAL_N_63_tmpxft_000019f5_00000000_10_DilatedMaxPool2d
ol_backward_nchw<float, float>(int, float const *, long const *, int, int, int, int, int, int, int, int, at::native::GLOBAL
_DilatedMaxPool2d_compute_75_cpp1_i1_db999de0::max_pool_backward_nchw<float, float>*)
6.22% 308.05ms 542 568.37us 1.1200us 17.017ms [CUDA memcpy HtoD]
5.33% 263.65ms 406 649.39us 1.2480us 87.996ms [CUDA memcpy DtoH]
4.18% 206.85ms 783 264.18us 127.84us 496.25us sgemm_32x32x32_NT_vec
4.15% 205.56ms 4564 45.039us 1.3110us 572.06us _ZN2at6native29vectorized_elementwise_kernelILi4EZZNS0_23gpu_kernel_with_scala
ensorIteratorEENKulvE_cIEvEUKulvE2_cIEvEULffe_BEvS4_RKT_EULfE0_NS_6detail5ArrayIPcLi2EEEEvTi0_T1
2.96% 146.61ms 663 221.13us 60.480us 431.20us sgemm_32x32x32_NT_vec
2.55% 126.17ms 263 479.74us 151.20us 545.05us maxwell_scudnn_128x64_relu_large_nn
2.48% 122.92ms 663 185.40us 43.680us 377.57us sgemm_128x128x8_TN_vec
```

As can be observed from the screenshot, we have used nvprof – a profiling tool to collect data about the GPU activity.

```
real    0m54.131s
user    0m42.959s
sys     0m8.413s
```

When we compared the runtime of the same job on the BareMetal and Cloud Environments, we noticed that the job ran much faster on the bare-metal system( as the time-elapsd was relatively lesser) than the cloud environment. The relative GPU time activity was the same. The throughput or the number of instructions processed per second is also relatively greater in the BareMetal environment.

### Example Run 2:

- **Running on Cloud:** For the second iteration, we train the model with the following hyperparameter settings: Batch size - 15 Epoch -3 and the Learning Rate has been set to 0.05. It



takes approximately 294 iterations and processes images in batches of 15 and passes over the entire dataset three times. We see an accuracy of about 25% at the end of the first epoch.

Training the model with the following hyperparameter settings: Batch size - 15 Epoch -3 and Learning Rate set to 0.05

```
!time nvprof python main.py -a alexnet -b 15 --epochs 3 --lr 0.05 Imagenet

==889== NVPROF is profiling process 889, command: python3 main.py -a alexnet -b 15 --epochs 3 --lr 0.05 Imagenet
=> creating model 'alexnet'
Epoch: [0] [ 0/294] Time 1.002 ( 1.002) Data 0.741 ( 0.741) Loss 6.9075e+00 (6.9075e+00) Acc@1 0.00 ( 0.00) Acc@5 0.00 ( 0.00)
Epoch: [0] [10/294] Time 0.036 ( 0.148) Data 0.004 ( 0.102) Loss 1.2042e+03 (1.2156e+02) Acc@1 20.00 ( 23.03) Acc@5 86.67 ( 75.15)
Epoch: [0] [20/294] Time 0.037 ( 0.144) Data 0.000 ( 0.106) Loss nan (nan) Acc@1 26.67 ( 23.81) Acc@5 100.00 ( 83.17)
Epoch: [0] [30/294] Time 0.032 ( 0.124) Data 0.000 ( 0.089) Loss nan (nan) Acc@1 33.33 ( 24.09) Acc@5 100.00 ( 88.60)
Epoch: [0] [40/294] Time 0.031 ( 0.124) Data 0.000 ( 0.091) Loss nan (nan) Acc@1 26.67 ( 25.69) Acc@5 100.00 ( 91.38)
Epoch: [0] [50/294] Time 0.034 ( 0.115) Data 0.000 ( 0.084) Loss nan (nan) Acc@1 13.33 ( 24.97) Acc@5 100.00 ( 93.07)
Epoch: [0] [60/294] Time 0.045 ( 0.114) Data 0.000 ( 0.084) Loss nan (nan) Acc@1 26.67 ( 24.70) Acc@5 100.00 ( 94.21)
Epoch: [0] [70/294] Time 0.028 ( 0.111) Data 0.000 ( 0.081) Loss nan (nan) Acc@1 13.33 ( 25.16) Acc@5 100.00 ( 95.02)
Epoch: [0] [80/294] Time 0.029 ( 0.111) Data 0.000 ( 0.081) Loss nan (nan) Acc@1 13.33 ( 25.10) Acc@5 100.00 ( 95.64)
Epoch: [0] [90/294] Time 0.032 ( 0.108) Data 0.000 ( 0.079) Loss nan (nan) Acc@1 40.00 ( 25.27) Acc@5 100.00 ( 96.12)
Epoch: [0] [100/294] Time 0.038 ( 0.109) Data 0.000 ( 0.081) Loss nan (nan) Acc@1 26.67 ( 25.54) Acc@5 100.00 ( 96.50)
Epoch: [0] [110/294] Time 0.032 ( 0.111) Data 0.005 ( 0.083) Loss nan (nan) Acc@1 33.33 ( 25.35) Acc@5 100.00 ( 96.82)
Epoch: [0] [120/294] Time 0.031 ( 0.110) Data 0.003 ( 0.083) Loss nan (nan) Acc@1 26.67 ( 25.73) Acc@5 100.00 ( 97.08)
Epoch: [0] [130/294] Time 0.025 ( 0.108) Data 0.000 ( 0.081) Loss nan (nan) Acc@1 46.67 ( 25.75) Acc@5 100.00 ( 97.30)
Epoch: [0] [140/294] Time 0.031 ( 0.109) Data 0.000 ( 0.082) Loss nan (nan) Acc@1 26.67 ( 25.72) Acc@5 100.00 ( 97.49)
Epoch: [0] [150/294] Time 0.029 ( 0.108) Data 0.000 ( 0.081) Loss nan (nan) Acc@1 20.00 ( 25.34) Acc@5 100.00 ( 97.66)
Epoch: [0] [160/294] Time 0.134 ( 0.109) Data 0.123 ( 0.082) Loss nan (nan) Acc@1 26.67 ( 25.38) Acc@5 100.00 ( 97.81)
Epoch: [0] [170/294] Time 0.024 ( 0.109) Data 0.000 ( 0.083) Loss nan (nan) Acc@1 20.00 ( 25.15) Acc@5 100.00 ( 97.93)
Epoch: [0] [180/294] Time 0.026 ( 0.109) Data 0.016 ( 0.083) Loss nan (nan) Acc@1 13.33 ( 24.79) Acc@5 100.00 ( 98.05)
Epoch: [0] [190/294] Time 0.035 ( 0.110) Data 0.000 ( 0.084) Loss nan (nan) Acc@1 13.33 ( 24.82) Acc@5 100.00 ( 98.15)
Epoch: [0] [200/294] Time 0.029 ( 0.109) Data 0.000 ( 0.083) Loss nan (nan) Acc@1 26.67 ( 25.14) Acc@5 100.00 ( 98.24)
Epoch: [0] [210/294] Time 0.161 ( 0.110) Data 0.149 ( 0.084) Loss nan (nan) Acc@1 6.67 ( 24.83) Acc@5 100.00 ( 98.33)
Epoch: [0] [220/294] Time 0.030 ( 0.109) Data 0.000 ( 0.083) Loss nan (nan) Acc@1 20.00 ( 24.86) Acc@5 100.00 ( 98.40)
Epoch: [0] [230/294] Time 0.072 ( 0.109) Data 0.055 ( 0.083) Loss nan (nan) Acc@1 26.67 ( 24.73) Acc@5 100.00 ( 98.47)
Epoch: [0] [240/294] Time 0.025 ( 0.108) Data 0.002 ( 0.082) Loss nan (nan) Acc@1 40.00 ( 24.76) Acc@5 100.00 ( 98.53)
Epoch: [0] [250/294] Time 0.034 ( 0.108) Data 0.000 ( 0.083) Loss nan (nan) Acc@1 20.00 ( 24.97) Acc@5 100.00 ( 98.59)
Epoch: [0] [260/294] Time 0.028 ( 0.108) Data 0.000 ( 0.083) Loss nan (nan) Acc@1 26.67 ( 25.08) Acc@5 100.00 ( 98.65)
Epoch: [0] [270/294] Time 0.042 ( 0.109) Data 0.003 ( 0.083) Loss nan (nan) Acc@1 20.00 ( 24.99) Acc@5 100.00 ( 98.70)
Epoch: [0] [280/294] Time 0.032 ( 0.108) Data 0.002 ( 0.083) Loss nan (nan) Acc@1 6.67 ( 24.93) Acc@5 100.00 ( 98.74)
Epoch: [0] [290/294] Time 0.028 ( 0.108) Data 0.000 ( 0.082) Loss nan (nan) Acc@1 26.67 ( 25.06) Acc@5 100.00 ( 98.79)
Test: [ 0/54] Time 1.077 ( 1.077) Loss nan (nan) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [10/54] Time 0.017 ( 0.171) Loss nan (nan) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [20/54] Time 0.481 ( 0.164) Loss nan (nan) Acc@1 100.00 ( 36.51) Acc@5 100.00 (100.00)
Test: [30/54] Time 0.011 ( 0.146) Loss nan (nan) Acc@1 0.00 ( 45.16) Acc@5 100.00 (100.00)
Test: [40/54] Time 0.261 ( 0.139) Loss nan (nan) Acc@1 0.00 ( 34.15) Acc@5 100.00 (100.00)
Test: [50/54] Time 0.009 ( 0.129) Loss nan (nan) Acc@1 0.00 ( 27.45) Acc@5 100.00 (100.00)
* Acc@1 25.926 Acc@5 100.000
Epoch: [1] [ 0/294] Time 0.707 ( 0.707) Data 0.681 ( 0.681) Loss nan (nan) Acc@1 26.67 ( 26.67) Acc@5 100.00 (100.00)
Epoch: [1] [10/294] Time 0.041 ( 0.157) Data 0.000 ( 0.131) Loss nan (nan) Acc@1 13.33 ( 22.42) Acc@5 100.00 (100.00)
Epoch: [1] [20/294] Time 0.181 ( 0.130) Data 0.166 ( 0.105) Loss nan (nan) Acc@1 20.00 ( 22.54) Acc@5 100.00 (100.00)
Epoch: [1] [30/294] Time 0.041 ( 0.125) Data 0.003 ( 0.098) Loss nan (nan) Acc@1 26.67 ( 23.44) Acc@5 100.00 (100.00)
Epoch: [1] [40/294] Time 0.038 ( 0.116) Data 0.002 ( 0.088) Loss nan (nan) Acc@1 20.00 ( 22.76) Acc@5 100.00 (100.00)
Epoch: [1] [50/294] Time 0.047 ( 0.125) Data 0.000 ( 0.095) Loss nan (nan) Acc@1 13.33 ( 23.14) Acc@5 100.00 (100.00)
```

Training the model with the following hyperparameter settings: Batch size - 15 Epoch -3 and Learning Rate set to 0.05

```
!time nvprof python main.py -a alexnet -b 15 --epochs 3 --lr 0.05 Imagenet > log2.txt

0.01% 1.5446ms 899 1.7180us 1.4400us 3.5200us cudnn::gemm::computeWgradBOffsetsKernel(cudnn::gemm::ComputeBOffsetsParams)
0.00% 1.1036ms 12 91.967us 26.399us 167.23us void fft2d_r2c_32x32<float, bool=0, unsigned int=0, bool=0>(float2*, float const *,
0.00% 817.83us 24 34.076us 23.263us 52.350us void cudnn::winoograd::generateWinoGradFilesKernel(int=0, float, float>(cudnn::wino
0.00% 719.47us 2 359.73us 321.11us 398.36us void fft2d_r2c_32x32<float, bool=0, unsigned int=5, bool=1>(float2*, float const *,
0.00% 676.11us 8 84.513us 22.591us 175.20us void fft2d_c2r_32x32<float, bool=0, bool=0, unsigned int=0, bool=0, bool=0>(float*
0.00% 62.877us 24 2.6190us 1.7920us 3.9680us compute_gemm_pointers(float2**, float2 const *, int, float2 const *, int, float2 c
0.00% 23.199us 12 1.9330us 1.0240us 2.9120us _ZN2at6native29vectorized_elementwise_kernelIli4EZN50_23gpu_kernel_with_scalarsIZZ
0.00% 16.703us 8 2.0870us 1.6320us 2.7830us cudnn::gemm::computeBOffsetsKernel(cudnn::gemm::ComputeBOffsetsParams)
0.00% 10.720us 4 2.6800us 2.4000us 2.9760us _ZN2at6native29vectorized_elementwise_kernelIli4EZZZN50_14gt_kernel_cudaERN5_14Te
0.00% 2.9120us 1 2.9120us 2.9120us 2.9120us _ZN2at6native29vectorized_elementwise_kernelIli4EZN50_23gpu_kernel_with_scalarsIZZ
0.00% 2.7830us 1 2.7830us 2.7830us 2.7830us _ZN2at6native29vectorized_elementwise_kernelIli4EZN50_23gpu_kernel_with_scalarsIZZ
API calls: 38.55% 9.26627s 3608 2.5683ms 9.2170us 33.230ms cudaMemcpyAsync
29.06% 6.98680s 42 166.35ms 9.9240us 6.97204s cudaMalloc
17.89% 4.30064s 208129 20.663us 5.3310us 14.274ms cudaLaunchKernel
6.14% 1.47622s 1307994 1.1280us 263ns 13.150ms cudaGetDevice
2.38% 571.15ms 103 5.5451ms 18.416us 99.746ms cudaMemcpy
1.29% 309.11ms 180 1.7173ms 385.90us 11.224ms cudaEventSynchronize
1.14% 273.36ms 4279 63.884us 534ns 14.517ms cudaEventDestroy
0.86% 205.94ms 31555 6.5260us 675ns 11.573ms cudaEventQuery
0.60% 144.20ms 230039 626ns 112ns 10.346ms cudaGetLastError
0.48% 115.86ms 6708 17.271us 4.5420us 5.8033ms cudaMemsetAsync
0.37% 89.008ms 22822 3.9000us 538ns 4.8319ms cudaEventRecord
0.36% 85.452ms 4320 19.780us 583ns 12.782ms cudaEventCreateWithFlags
0.23% 54.876ms 18 3.0487ms 38.003us 9.5926ms cudaHostAlloc
0.21% 50.665ms 1504 33.686us 2.2980us 8.0673ms cudaStreamSynchronize
0.13% 30.761ms 2088 14.732us 6.8830us 2.5510ms cudaPointerGetAttributes
0.10% 23.058ms 8424 2.7370us 906ns 3.1970ms cudaOccupancyMaxActiveBlocksPerMultiprocessorWithFlags
0.06% 14.109ms 1388 10.164us 594ns 7.2974ms cudaStreamWaitEvent
0.04% 10.550ms 6272 1.6820us 320ns 3.1825ms cudaSetDevice
0.03% 6.5844ms 31 212.40us 832ns 1.1747ms cudaFree
0.03% 6.4983ms 6336 1.0250us 110ns 292.18us cudaGetDeviceCount
0.02% 4.5162ms 676 6.6800us 651ns 3.4846ms cudaFuncSetAttribute
0.02% 3.6157ms 2104 1.7180us 298ns 772.39us cuDevicePrimaryCtxGetState
0.01% 3.1897ms 95 33.576us 2.5390us 897.94us cudaBindTexture
0.01% 1.9208ms 28 68.601us 20.669us 930.65us cudaMemGetInfo
0.01% 1.3567ms 72 18.842us 2.1930us 271.46us cudaStreamCreateWithPriority
0.00% 1.1246ms 180 6.2470us 3.2520us 33.922us cudaEventElapsedTime
0.00% 733.57us 2 366.78us 347.11us 386.46us cuDeviceTotalMem
```



Once all three epochs' have been processed and the training is complete. The overall time taken to process has comparatively increased for this run as compared to the previous example -could be expected from the increase in epoch and decrease in batch size. The learning rate was increased and hence the step size has also increased. One can also note that the accuracy didn't change much.

Profiling the model with the following hyperparameter settings: Batch size - 15 Epoch -3 and Learning Rate set to 0.05

```
!time nvprof python main.py -a alexnet -b 15 --epochs 3 --lr 0.05 Imagenet
```

Test: [40/54]	Time 0.309 ( 0.139)	Loss nan (nan)	Acc@1 0.00 ( 34.15)	Acc@5 100.00 (100.00)			
Test: [50/54]	Time 0.317 ( 0.131)	Loss nan (nan)	Acc@1 0.00 ( 27.45)	Acc@5 100.00 (100.00)			
* Acc@1 25.926 Acc@5 100.000							
==889== Profiling application: python3 main.py -a alexnet -b 15 --epochs 3 --lr 0.05 Imagenet							
==889== Profiling result:							
Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	38.75%	10.4732s	58492	179.05us	1.0560us	2.0378ms	_ZN2at6native29vectorized_elementwise_kernelILi4EZZN50_15add_kernel
	9.31%	2.51664s	8825	285.17us	77.312us	939.97us	volta_sgemv_128x64_nt
	6.70%	1.81174s	18272	99.153us	1.0240us	1.4527ms	_ZN2at6native29vectorized_elementwise_kernelILi4EZZN50_23gpu_kerne
	5.60%	1.51325s	3132	483.16us	105.63us	1.2121ms	volta_sgemv_128x32_sliced1x4_tn
	5.19%	1.40295s	5062	277.15us	166.31us	724.32us	volta_sgemv_128x64_nn
	3.75%	1.01387s	18506	54.786us	864ns	653.60us	_ZN2at6native29vectorized_elementwise_kernelILi4EZZN50_16fill_ke
	3.47%	937.20ms	2646	354.20us	80.704us	1.1238ms	volta_sgemv_128x32_sliced1x4_nn
	3.39%	916.01ms	2108	434.54us	1.0560us	33.822ms	[CUDA memcpy HtoD]
	2.19%	590.98ms	883	669.28us	537.35us	1.4363ms	volta_scudnn_128x64_strided8_splitK_medium_nn_v1
	2.02%	547.24ms	1587	344.83us	1.2160us	101.37ms	[CUDA memcpy DtoH]
	1.97%	533.42ms	13482	39.565us	1.9210us	152.10us	_ZN2at6native29vectorized_elementwise_kernelILi4EZZN50_21threshol
	1.79%	484.31ms	1046	463.01us	276.07us	874.72us	volta_scudnn_128x64_relu_xregs_large_nn_v1
	1.64%	442.45ms	2646	167.21us	47.552us	466.82us	void at::native::GLOBAL_N_63_tmpxft_00002580_00000000_10_Dilat
	1.61%	435.10ms	5775	75.342us	58.529us	113.12us	void cudnn::winograd_nonfused::winogradForwardFilter4x4x4<float, f
	1.45%	391.41ms	5775	67.777us	37.600us	116.51us	void cudnn::winograd_nonfused::winogradForwardData4x4x4<float, floa
	1.43%	385.85ms	5775	66.813us	41.696us	98.529us	void cudnn::winograd_nonfused::winogradForwardOutput4x4x4<float, f
	1.07%	289.17ms	2814	102.76us	50.176us	267.14us	void cudnn::winograd_nonfused::winogradForwardData9x9_5x5<float, f
	1.00%	269.49ms	2652	101.62us	73.376us	237.28us	void cudnn::winograd_nonfused::winogradWgradOutput4x4x4<float, floa
	0.99%	266.36ms	5220	51.026us	16.384us	126.69us	_ZN2at6native27unrolled_elementwise_kernelILi4EZZN50_15add_kernel_cl
	0.83%	223.28ms	1930	115.69us	51.840us	213.54us	void cudnn::winograd_nonfused::winogradForwardOutput9x9_5x5<float
	0.71%	193.20ms	3132	61.687us	9.4720us	120.51us	void at::native::GLOBAL_N_63_tmpxft_00002580_00000000_10_Dilat
	0.69%	186.13ms	2652	70.185us	39.264us	101.31us	void cudnn::winograd_nonfused::winogradWgradDelta4x4x4<float, float
	0.67%	180.30ms	6498	27.747us	2.2080us	72.896us	_ZN2at6native13reduce_kernelILi512ELi1ENS0_8ReduceOpIFNS0_14func
	0.61%	164.24ms	2652	61.932us	21.888us	98.817us	void cudnn::winograd_nonfused::winogradWgradData4x4x4<float, float
	0.49%	133.71ms	884	151.25us	127.07us	207.23us	void cudnn::winograd_nonfused::winogradWgradDelta9x9_5x5<float, f
	0.45%	120.66ms	1930	62.520us	59.840us	71.872us	void cudnn::winograd_nonfused::winogradForwardFilter9x9_5x5<float
	0.28%	76.978ms	1044	73.733us	45.377us	138.79us	void at::native::GLOBAL_N_69_tmpxft_00001db3_00000000_10_Adapt
	0.22%	59.995ms	884	67.868us	62.944us	98.432us	void cudnn::winograd_nonfused::winogradWgradOutput9x9_5x5<float,
	0.19%	52.005ms	882	58.962us	40.992us	139.62us	void at::native::GLOBAL_N_69_tmpxft_00001db3_00000000_10_Adapt
	0.13%	34.859ms	16	2.1787ms	337.31us	5.0659ms	volta_cgconv_32x32_nt
	0.12%	33.689ms	1044	32.269us	13.536us	56.833us	void gatherTopK<float, unsigned int, int=2, bool=1>(TensorInfo<f
	0.09%	24.820ms	16	1.5512ms	675.72us	3.0678ms	void cudnn::detail::dgrad_engine<float, int=128, int=6, int=8, in
	0.08%	22.195ms	2646	8.3880us	6.4000us	14.112us	_ZN2at6native13reduce_kernelILi128ELi4ENS0_8ReduceOpIFNS0_14func
	0.08%	21.230ms	3132	6.7780us	4.5760us	11.840us	_ZN2at6native27unrolled_elementwise_kernelILi4EZZN50_21copy_device_f
	0.07%	19.008ms	10	1.9008ms	829.12us	4.2025ms	void cudnn::detail::implicit_convolve_sgemv<float, float, int=512
	0.06%	17.255ms	24	718.98us	310.18us	1.3707ms	volta_cgconv_32x32_tn
	0.06%	17.097ms	12	1.4247ms	523.52us	2.5385ms	void cudnn::detail::wgrad_alg0_engine<float, int=128, int=6, int
	0.06%	14.952ms	27	553.78us	325.73us	1.4900ms	volta_scudnn_winograd_128x128_ldg1_ldg4_relu_tile148t_nt_v1
	0.05%	14.027ms	6	2.3379ms	1.7274ms	3.4258ms	void fft2d_r2c_32x32<float, bool=0, unsigned int=1, bool=1>(float
	0.05%	13.336ms	18	740.89us	86.625us	2.3339ms	void fft2d_r2c_32x32<float, bool=0, unsigned int=1, bool=0>(float
	0.04%	9.6504ms	1764	5.4700us	3.6480us	14.368us	_ZN2at6native29vectorized_elementwise_kernelILi4EZZN50_78_GLOBAL

real	2m24.908s
user	3m15.853s
sys	0m33.090s

- **Running on BareMetal:** For the second iteration, we train the model with the same hyperparameter settings: Batch size - 15 Epoch -3 and the Learning Rate has been set to 0.05. It takes approximately 294 iterations and processes images in batches of 15 and passes over the entire dataset three times.

```
(py3.6.3) [as13594@gpu-61 ~]$ time nvprof python /home/as13594/examples/imagenet/main.py -a alexnet -b 15 --epochs 3 --lr 0.05 /home/as13594/
==46536== NVPROF is profiling process 46536, command: python /home/as13594/examples/imagenet/main.py -a alexnet -b 15 --epochs 3 --lr 0.05 /home/as13594/
-> creating model 'alexnet'
Epoch: [0] [ 0/294] Time 1.047 ( 1.047) Data 0.814 ( 0.814) Loss 6.9044e+00 (6.9044e+00) Acc@1 0.00 ( 0.00) Acc@5 0.00 ( 0.00)
Epoch: [0] [10/294] Time 0.312 ( 0.202) Data 0.303 ( 0.158) Loss 6.2731e+00 (6.4928e+00) Acc@1 33.33 (21.21) Acc@5 100.00 ( 83.03)
Epoch: [0] [20/294] Time 0.031 ( 0.169) Data 0.000 ( 0.134) Loss nan (nan) Acc@1 13.33 (22.22) Acc@5 100.00 ( 77.46)
Epoch: [0] [30/294] Time 0.308 ( 0.153) Data 0.296 ( 0.121) Loss nan (nan) Acc@1 6.67 (20.65) Acc@5 100.00 ( 84.73)
Epoch: [0] [40/294] Time 0.043 ( 0.138) Data 0.010 ( 0.106) Loss nan (nan) Acc@1 26.67 (20.65) Acc@5 100.00 ( 88.46)
Epoch: [0] [50/294] Time 0.315 ( 0.134) Data 0.306 ( 0.105) Loss nan (nan) Acc@1 26.67 (21.70) Acc@5 100.00 ( 90.72)
Epoch: [0] [60/294] Time 0.037 ( 0.128) Data 0.001 ( 0.099) Loss nan (nan) Acc@1 20.00 (22.08) Acc@5 100.00 ( 92.24)
Epoch: [0] [70/294] Time 0.334 ( 0.128) Data 0.325 ( 0.099) Loss nan (nan) Acc@1 40.00 (22.72) Acc@5 100.00 ( 93.33)
Epoch: [0] [80/294] Time 0.044 ( 0.124) Data 0.002 ( 0.094) Loss nan (nan) Acc@1 20.00 (23.70) Acc@5 100.00 ( 94.16)
Epoch: [0] [90/294] Time 0.580 ( 0.126) Data 0.568 ( 0.097) Loss nan (nan) Acc@1 20.00 (23.30) Acc@5 100.00 ( 94.80)
Epoch: [0] [100/294] Time 0.042 ( 0.123) Data 0.012 ( 0.094) Loss nan (nan) Acc@1 20.00 (23.04) Acc@5 100.00 ( 95.31)
Epoch: [0] [110/294] Time 0.419 ( 0.123) Data 0.409 ( 0.095) Loss nan (nan) Acc@1 13.33 (23.24) Acc@5 100.00 ( 95.74)
Epoch: [0] [120/294] Time 0.041 ( 0.120) Data 0.013 ( 0.092) Loss nan (nan) Acc@1 20.00 (22.87) Acc@5 100.00 ( 96.09)
Epoch: [0] [130/294] Time 0.282 ( 0.121) Data 0.273 ( 0.092) Loss nan (nan) Acc@1 46.67 (23.82) Acc@5 100.00 ( 96.39)
Epoch: [0] [140/294] Time 0.036 ( 0.119) Data 0.000 ( 0.091) Loss nan (nan) Acc@1 26.67 (23.88) Acc@5 100.00 ( 96.64)
Epoch: [0] [150/294] Time 0.498 ( 0.120) Data 0.487 ( 0.092) Loss nan (nan) Acc@1 13.33 (23.75) Acc@5 100.00 ( 96.87)
Epoch: [0] [160/294] Time 0.041 ( 0.118) Data 0.002 ( 0.091) Loss nan (nan) Acc@1 20.00 (23.77) Acc@5 100.00 ( 97.06)
Epoch: [0] [170/294] Time 0.300 ( 0.120) Data 0.291 ( 0.092) Loss nan (nan) Acc@1 33.33 (24.76) Acc@5 100.00 ( 97.23)
Epoch: [0] [180/294] Time 0.033 ( 0.118) Data 0.004 ( 0.091) Loss nan (nan) Acc@1 33.33 (24.42) Acc@5 100.00 ( 97.38)
Epoch: [0] [190/294] Time 0.314 ( 0.118) Data 0.305 ( 0.091) Loss nan (nan) Acc@1 20.00 (24.75) Acc@5 100.00 ( 97.52)
Epoch: [0] [200/294] Time 0.034 ( 0.117) Data 0.000 ( 0.089) Loss nan (nan) Acc@1 26.67 (24.34) Acc@5 100.00 ( 97.65)
Epoch: [0] [210/294] Time 0.337 ( 0.118) Data 0.328 ( 0.090) Loss nan (nan) Acc@1 40.00 (24.49) Acc@5 100.00 ( 97.76)
Epoch: [0] [220/294] Time 0.035 ( 0.117) Data 0.002 ( 0.089) Loss nan (nan) Acc@1 26.67 (24.68) Acc@5 100.00 ( 97.86)
Epoch: [0] [230/294] Time 0.319 ( 0.117) Data 0.308 ( 0.090) Loss nan (nan) Acc@1 33.33 (24.76) Acc@5 100.00 ( 97.95)
Epoch: [0] [240/294] Time 0.032 ( 0.116) Data 0.000 ( 0.089) Loss nan (nan) Acc@1 26.67 (24.67) Acc@5 100.00 ( 98.04)
Epoch: [0] [250/294] Time 0.232 ( 0.116) Data 0.222 ( 0.089) Loss nan (nan) Acc@1 6.67 (24.46) Acc@5 100.00 ( 98.11)
Epoch: [0] [260/294] Time 0.030 ( 0.115) Data 0.000 ( 0.089) Loss nan (nan) Acc@1 26.67 (24.44) Acc@5 100.00 ( 98.19)
Epoch: [0] [270/294] Time 0.031 ( 0.115) Data 0.005 ( 0.089) Loss nan (nan) Acc@1 26.67 (24.43) Acc@5 100.00 ( 98.25)
Epoch: [0] [280/294] Time 0.032 ( 0.116) Data 0.000 ( 0.090) Loss nan (nan) Acc@1 26.67 (24.32) Acc@5 100.00 ( 98.32)
Epoch: [0] [290/294] Time 0.026 ( 0.115) Data 0.000 ( 0.089) Loss nan (nan) Acc@1 13.33 (24.56) Acc@5 100.00 ( 98.37)
Test: [ 0/54] Time 1.056 ( 1.056) Loss nan (nan) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [10/54] Time 0.004 ( 0.164) Loss nan (nan) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [20/54] Time 0.448 ( 0.169) Loss nan (nan) Acc@1 100.00 (36.51) Acc@5 100.00 (100.00)
Test: [30/54] Time 0.004 ( 0.144) Loss nan (nan) Acc@1 0.00 (43.01) Acc@5 100.00 (100.00)
Test: [40/54] Time 0.433 ( 0.141) Loss nan (nan) Acc@1 0.00 (32.52) Acc@5 100.00 (100.00)
Test: [50/54] Time 0.004 ( 0.127) Loss nan (nan) Acc@1 0.00 (26.14) Acc@5 100.00 (100.00)
* Acc@1 25.000 Acc@5 100.000
Epoch: [1] [ 0/294] Time 0.647 ( 0.647) Data 0.637 ( 0.637) Loss nan (nan) Acc@1 33.33 (33.33) Acc@5 100.00 (100.00)
Epoch: [1] [10/294] Time 0.024 ( 0.141) Data 0.000 ( 0.120) Loss nan (nan) Acc@1 13.33 (26.06) Acc@5 100.00 (100.00)
Epoch: [1] [20/294] Time 0.225 ( 0.129) Data 0.216 ( 0.110) Loss nan (nan) Acc@1 20.00 (25.08) Acc@5 100.00 (100.00)
Epoch: [1] [30/294] Time 0.036 ( 0.119) Data 0.000 ( 0.100) Loss nan (nan) Acc@1 6.67 (25.16) Acc@5 100.00 (100.00)
Epoch: [1] [40/294] Time 0.257 ( 0.127) Data 0.248 ( 0.107) Loss nan (nan) Acc@1 20.00 (25.20) Acc@5 100.00 (100.00)
Epoch: [1] [50/294] Time 0.034 ( 0.129) Data 0.000 ( 0.107) Loss nan (nan) Acc@1 20.00 (24.97) Acc@5 100.00 (100.00)
Epoch: [1] [60/294] Time 0.268 ( 0.127) Data 0.258 ( 0.104) Loss nan (nan) Acc@1 20.00 (23.50) Acc@5 100.00 (100.00)
Epoch: [1] [70/294] Time 0.030 ( 0.122) Data 0.000 ( 0.098) Loss nan (nan) Acc@1 13.33 (23.38) Acc@5 100.00 (100.00)
Epoch: [1] [80/294] Time 0.257 ( 0.121) Data 0.287 ( 0.097) Loss nan (nan) Acc@1 20.00 (23.37) Acc@5 100.00 (100.00)
Epoch: [1] [90/294] Time 0.046 ( 0.118) Data 0.007 ( 0.093) Loss nan (nan) Acc@1 26.67 (23.30) Acc@5 100.00 (100.00)
Epoch: [1] [100/294] Time 0.221 ( 0.118) Data 0.212 ( 0.093) Loss nan (nan) Acc@1 20.00 (23.10) Acc@5 100.00 (100.00)
```

As can be observed in the screenshot, the profiling information for the run has been generated at the end of the third epoch by nvprof.

```
Epoch: [2] [190/294] Time 0.039 ( 0.113) Data 0.000 ( 0.088) Loss nan (nan) Acc@1 13.33 (24.92) Acc@5 100.00 (100.00)
Epoch: [2] [200/294] Time 0.284 ( 0.114) Data 0.275 ( 0.089) Loss nan (nan) Acc@1 13.33 (24.64) Acc@5 100.00 (100.00)
Epoch: [2] [210/294] Time 0.045 ( 0.113) Data 0.000 ( 0.088) Loss nan (nan) Acc@1 20.00 (24.93) Acc@5 100.00 (100.00)
Epoch: [2] [220/294] Time 0.327 ( 0.114) Data 0.318 ( 0.089) Loss nan (nan) Acc@1 26.67 (24.92) Acc@5 100.00 (100.00)
Epoch: [2] [230/294] Time 0.045 ( 0.113) Data 0.000 ( 0.088) Loss nan (nan) Acc@1 13.33 (24.85) Acc@5 100.00 (100.00)
Epoch: [2] [240/294] Time 0.246 ( 0.113) Data 0.236 ( 0.088) Loss nan (nan) Acc@1 33.33 (24.76) Acc@5 100.00 (100.00)
Epoch: [2] [250/294] Time 0.026 ( 0.112) Data 0.000 ( 0.088) Loss nan (nan) Acc@1 26.67 (24.78) Acc@5 100.00 (100.00)
Epoch: [2] [260/294] Time 0.046 ( 0.112) Data 0.037 ( 0.088) Loss nan (nan) Acc@1 40.00 (24.93) Acc@5 100.00 (100.00)
Epoch: [2] [270/294] Time 0.026 ( 0.113) Data 0.000 ( 0.089) Loss nan (nan) Acc@1 20.00 (25.04) Acc@5 100.00 (100.00)
Epoch: [2] [280/294] Time 0.193 ( 0.113) Data 0.184 ( 0.088) Loss nan (nan) Acc@1 0.00 (25.01) Acc@5 100.00 (100.00)
Epoch: [2] [290/294] Time 0.026 ( 0.112) Data 0.000 ( 0.088) Loss nan (nan) Acc@1 26.67 (24.90) Acc@5 100.00 (100.00)
Test: [ 0/54] Time 1.070 ( 1.070) Loss nan (nan) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [10/54] Time 0.004 ( 0.165) Loss nan (nan) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [20/54] Time 0.408 ( 0.166) Loss nan (nan) Acc@1 100.00 (36.51) Acc@5 100.00 (100.00)
Test: [30/54] Time 0.005 ( 0.141) Loss nan (nan) Acc@1 0.00 (43.01) Acc@5 100.00 (100.00)
Test: [40/54] Time 0.418 ( 0.138) Loss nan (nan) Acc@1 0.00 (32.52) Acc@5 100.00 (100.00)
Test: [50/54] Time 0.005 ( 0.124) Loss nan (nan) Acc@1 0.00 (26.14) Acc@5 100.00 (100.00)
* Acc@1 25.000 Acc@5 100.000
==46536== Profiling application: python /home/as13594/examples/imagenet/main.py -a alexnet -b 15 --epochs 3 --lr 0.05 /home/as13594/
==46536== Profiling result:
Type Time(%) Time Calls Avg Min Max Name
GPU activities: 29.41% 4.69109s 58492 80.200us 1.2800us 862.46us _ZN2at6native29vectorized_elementwise_kernelIli4EZZZ2NS0_15a
lve_cIEVENKulvE2_cIEVEUllfE_N3_gdetail5ArrayIPcLi3EEEEvIT0_T1_
7.71% 1.22993s 6185 198.86us 121.70us 269.53us maxwell_sgemv_128x64_nt
6.37% 1.01638s 2646 384.12us 101.73us 600.48us void at::native::GLOBAL_N_63_tmpxft_000019f5_00000000_10
ol_backward_nchw<float, float>(int, float const *, long const *, int, int, int, int, int, int, int, int, int, int, int, int, at
_DilatedMaxPool2d_compute_75_cppl_i_db999de0::max_pool_backward_nchw(float, float*)
6.07% 968.75ms 5057 191.57us 128.35us 289.37us maxwell_sgemv_128x64_nn
5.27% 841.12ms 3132 268.56us 127.07us 496.09us sgemm_32x32x32_NT_vec
5.16% 822.56ms 18272 45.017us 1.3110us 591.20us _ZN2at6native29vectorized_elementwise_kernelIli4EZZNS0_23gpn
ensorIteratorEENKulvE2_cIEVEUllfE_EEVS4_RKT_EUllfE0_N3_gdetail5ArrayIPcLi2EEEEvIT0_T1_
5.10% 813.42ms 2108 385.87us 1.1200us 17.000ms [CUDA memcpy HtoD]
4.77% 760.06ms 1587 478.93us 1.2480us 87.019ms [CUDA memcpy DtoH]
3.76% 600.10ms 2646 226.79us 61.120us 435.55us sgemm_32x32x32_NN_vec
2.76% 440.57ms 2646 166.51us 45.823us 322.33us sgemm_128x128x8_TN_vec
2.67% 425.72ms 18506 23.004us 1.0880us 275.97us _ZN2at6native29vectorized_elementwise_kernelIli4EZZZ2NS0_16f
UlvE_cIEVENKulvE2_cIEVEUllvE_N3_gdetail5ArrayIPcLi1EEEEvIT0_T1_
2.42% 386.45ms 1047 369.11us 152.74us 419.45us maxwell_scudnn_128x64_relu_large_nn
2.11% 336.01ms 886 379.24us 254.14us 833.28us maxwell_scudnn_128x64_stridedB_splitK_medium_nn
```

```
real 2m33.502s
user 2m7.061s
sys 0m20.072s
```

When we compare the runtime of the same job on the BareMetal and Cloud Environments, we noticed that the time elapsed in both cases was relatively same. The relative GPU time activity on Cloud was comparatively more.

## Example Run 3:

**Running on Cloud:** We train the model with the following hyperparameter settings: Batch size – 10 Epoch -5 and the Learning Rate has been set to 0.001. It takes approximately 441 iterations in each epoch and processes images in batches of 10.

Profiling the model with the following hyperparameter settings: Batch size – 10 Epoch -5 and the Learning Rate set to 0.001.

```
!time python main.py -a alexnet -b 10 --epochs 5 --lr 0.001 Imagenet

=> creating model 'alexnet'
Epoch: [0][ 0/441] Time 0.813 ( 0.813) Data 0.563 ( 0.563) Loss 6.8972e+00 (6.8972e+00) Acc@1 0.00 ( 0.00) Acc@5 0.00 ( 0.00)
Epoch: [0][10/441] Time 0.256 ( 0.130) Data 0.247 ( 0.086) Loss 6.8794e+00 (6.8933e+00) Acc@1 20.00 ( 6.36) Acc@5 30.00 (10.00)
Epoch: [0][20/441] Time 0.028 ( 0.095) Data 0.002 ( 0.061) Loss 6.8439e+00 (6.8784e+00) Acc@1 30.00 (17.14) Acc@5 100.00 (44.29)
Epoch: [0][30/441] Time 0.177 ( 0.087) Data 0.169 ( 0.057) Loss 6.7835e+00 (6.8592e+00) Acc@1 40.00 (18.71) Acc@5 100.00 (62.26)
Epoch: [0][40/441] Time 0.035 ( 0.084) Data 0.000 ( 0.056) Loss 6.6883e+00 (6.8322e+00) Acc@1 40.00 (20.00) Acc@5 100.00 (71.46)
Epoch: [0][50/441] Time 0.116 ( 0.080) Data 0.108 ( 0.053) Loss 6.6436e+00 (6.7841e+00) Acc@1 10.00 (20.59) Acc@5 100.00 (77.06)
Epoch: [0][60/441] Time 0.029 ( 0.076) Data 0.002 ( 0.050) Loss 6.4384e+00 (6.4995e+00) Acc@1 20.00 (20.33) Acc@5 100.00 (80.82)
Epoch: [0][70/441] Time 0.117 ( 0.074) Data 0.109 ( 0.049) Loss 6.3897e+00 (6.0888e+00) Acc@1 30.00 (20.00) Acc@5 100.00 (83.52)
Epoch: [0][80/441] Time 0.063 ( 0.073) Data 0.055 ( 0.048) Loss 6.2438e+00 (5.7899e+00) Acc@1 50.00 (21.48) Acc@5 100.00 (84.94)
Epoch: [0][90/441] Time 0.066 ( 0.071) Data 0.057 ( 0.047) Loss 6.2172e+00 (5.5987e+00) Acc@1 40.00 (21.32) Acc@5 100.00 (85.49)
Epoch: [0][100/441] Time 0.031 ( 0.072) Data 0.005 ( 0.048) Loss 6.1718e+00 (5.2987e+00) Acc@1 30.00 (21.58) Acc@5 100.00 (86.93)
Epoch: [0][110/441] Time 0.032 ( 0.072) Data 0.000 ( 0.048) Loss 6.0917e+00 (5.0044e+00) Acc@1 20.00 (21.62) Acc@5 100.00 (88.11)
Epoch: [0][120/441] Time 0.029 ( 0.071) Data 0.000 ( 0.047) Loss 6.1617e+00 (4.7425e+00) Acc@1 10.00 (21.90) Acc@5 100.00 (89.09)
Epoch: [0][130/441] Time 0.028 ( 0.071) Data 0.000 ( 0.047) Loss 6.4426e+00 (4.4971e+00) Acc@1 30.00 (22.67) Acc@5 100.00 (89.92)
Epoch: [0][140/441] Time 0.032 ( 0.072) Data 0.000 ( 0.048) Loss 6.1419e+00 (4.2848e+00) Acc@1 30.00 (22.98) Acc@5 100.00 (90.64)
Epoch: [0][150/441] Time 0.040 ( 0.071) Data 0.006 ( 0.048) Loss 6.4089e+00 (4.0995e+00) Acc@1 30.00 (23.05) Acc@5 100.00 (91.26)
Epoch: [0][160/441] Time 0.032 ( 0.071) Data 0.002 ( 0.047) Loss 6.4663e+00 (3.9421e+00) Acc@1 20.00 (22.86) Acc@5 100.00 (91.80)
Epoch: [0][170/441] Time 0.165 ( 0.071) Data 0.157 ( 0.048) Loss 6.4658e+00 (3.7951e+00) Acc@1 10.00 (22.81) Acc@5 100.00 (92.28)
Epoch: [0][180/441] Time 0.031 ( 0.070) Data 0.002 ( 0.047) Loss 6.1791e+00 (3.6673e+00) Acc@1 10.00 (22.87) Acc@5 100.00 (92.71)
Epoch: [0][190/441] Time 0.056 ( 0.070) Data 0.048 ( 0.048) Loss 6.4729e+00 (3.5537e+00) Acc@1 20.00 (23.09) Acc@5 100.00 (93.09)
Epoch: [0][200/441] Time 0.105 ( 0.070) Data 0.097 ( 0.048) Loss 6.1443e+00 (3.4520e+00) Acc@1 60.00 (23.43) Acc@5 100.00 (93.43)
Epoch: [0][210/441] Time 0.028 ( 0.070) Data 0.000 ( 0.048) Loss 6.6849e+00 (3.3612e+00) Acc@1 10.00 (23.46) Acc@5 100.00 (93.74)
Epoch: [0][220/441] Time 0.146 ( 0.070) Data 0.138 ( 0.048) Loss 6.5837e+00 (3.2794e+00) Acc@1 10.00 (23.17) Acc@5 100.00 (94.03)
Epoch: [0][230/441] Time 0.085 ( 0.070) Data 0.077 ( 0.048) Loss 6.1623e+00 (3.2018e+00) Acc@1 20.00 (23.38) Acc@5 100.00 (94.29)
Epoch: [0][240/441] Time 0.026 ( 0.070) Data 0.000 ( 0.048) Loss 6.6338e+00 (3.1287e+00) Acc@1 0.00 (23.44) Acc@5 100.00 (94.52)
Epoch: [0][250/441] Time 0.147 ( 0.070) Data 0.138 ( 0.048) Loss 6.5658e+00 (3.0687e+00) Acc@1 10.00 (23.11) Acc@5 100.00 (94.74)
Epoch: [0][260/441] Time 0.057 ( 0.069) Data 0.048 ( 0.048) Loss 6.5897e+00 (3.0108e+00) Acc@1 0.00 (22.99) Acc@5 100.00 (94.94)
Epoch: [0][270/441] Time 0.058 ( 0.069) Data 0.048 ( 0.048) Loss 6.5185e+00 (2.9545e+00) Acc@1 30.00 (23.06) Acc@5 100.00 (95.13)
Epoch: [0][280/441] Time 0.169 ( 0.069) Data 0.161 ( 0.048) Loss 6.6898e+00 (2.9011e+00) Acc@1 0.00 (23.20) Acc@5 100.00 (95.30)
Epoch: [0][290/441] Time 0.028 ( 0.069) Data 0.000 ( 0.048) Loss 6.3844e+00 (2.8514e+00) Acc@1 20.00 (23.47) Acc@5 100.00 (95.46)
Epoch: [0][300/441] Time 0.026 ( 0.070) Data 0.000 ( 0.049) Loss 6.4348e+00 (2.8068e+00) Acc@1 40.00 (23.49) Acc@5 100.00 (95.61)
Epoch: [0][310/441] Time 0.036 ( 0.070) Data 0.000 ( 0.048) Loss 6.5076e+00 (2.7652e+00) Acc@1 20.00 (23.44) Acc@5 100.00 (95.76)
Epoch: [0][320/441] Time 0.034 ( 0.069) Data 0.000 ( 0.048) Loss 6.4182e+00 (2.7256e+00) Acc@1 10.00 (23.33) Acc@5 100.00 (95.89)
Epoch: [0][330/441] Time 0.027 ( 0.069) Data 0.000 ( 0.048) Loss 6.4312e+00 (2.6859e+00) Acc@1 30.00 (23.41) Acc@5 100.00 (96.01)
Epoch: [0][340/441] Time 0.029 ( 0.070) Data 0.000 ( 0.048) Loss 6.3906e+00 (2.6504e+00) Acc@1 20.00 (23.34) Acc@5 100.00 (96.13)
Epoch: [0][350/441] Time 0.029 ( 0.070) Data 0.021 ( 0.048) Loss 6.2557e+00 (2.6155e+00) Acc@1 50.00 (23.56) Acc@5 100.00 (96.24)
Epoch: [0][360/441] Time 0.145 ( 0.070) Data 0.137 ( 0.048) Loss 6.2984e+00 (2.5825e+00) Acc@1 50.00 (23.57) Acc@5 100.00 (96.34)
Epoch: [0][370/441] Time 0.028 ( 0.070) Data 0.000 ( 0.048) Loss 6.5003e+00 (2.5506e+00) Acc@1 10.00 (23.67) Acc@5 100.00 (96.44)
Epoch: [0][380/441] Time 0.158 ( 0.070) Data 0.147 ( 0.048) Loss 6.2910e+00 (2.5215e+00) Acc@1 40.00 (23.81) Acc@5 100.00 (96.54)
Epoch: [0][390/441] Time 0.030 ( 0.069) Data 0.003 ( 0.048) Loss 6.1612e+00 (2.4964e+00) Acc@1 30.00 (23.94) Acc@5 100.00 (96.62)
Epoch: [0][400/441] Time 0.169 ( 0.069) Data 0.161 ( 0.048) Loss 6.4447e+00 (2.4688e+00) Acc@1 20.00 (24.29) Acc@5 100.00 (96.71)
Epoch: [0][410/441] Time 0.155 ( 0.070) Data 0.147 ( 0.048) Loss 6.3850e+00 (2.4426e+00) Acc@1 30.00 (24.43) Acc@5 100.00 (96.79)
Epoch: [0][420/441] Time 0.025 ( 0.069) Data 0.000 ( 0.048) Loss 6.3115e+00 (2.4172e+00) Acc@1 40.00 (24.68) Acc@5 100.00 (96.86)
Epoch: [0][430/441] Time 0.175 ( 0.069) Data 0.167 ( 0.048) Loss 6.3211e+00 (2.3923e+00) Acc@1 30.00 (24.85) Acc@5 100.00 (96.94)
Epoch: [0][440/441] Time 0.131 ( 0.069) Data 0.000 ( 0.048) Loss 6.1985e+00 (2.3684e+00) Acc@1 40.00 (25.18) Acc@5 100.00 (97.00)
Test: [ 0/81] Time 0.855 ( 0.855) Loss 1.5621e+00 (1.5621e+00) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [10/81] Time 0.006 ( 0.125) Loss 1.3568e+00 (1.3213e+00) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [20/81] Time 0.008 ( 0.101) Loss 1.3765e+00 (1.3235e+00) Acc@1 10.00 ( 0.48) Acc@5 100.00 (100.00)
Test: [30/81] Time 0.010 ( 0.089) Loss 1.2447e+00 (1.3082e+00) Acc@1 10.00 ( 6.45) Acc@5 100.00 (100.00)
Test: [40/81] Time 0.009 ( 0.088) Loss 1.2227e+00 (1.3024e+00) Acc@1 40.00 ( 8.78) Acc@5 100.00 (100.00)
Test: [50/81] Time 0.082 ( 0.084) Loss 8.5190e-01 (1.2285e+00) Acc@1 100.00 (25.69) Acc@5 100.00 (100.00)
Test: [60/81] Time 0.086 ( 0.083) Loss 9.7332e-01 (1.1812e+00) Acc@1 80.00 (36.72) Acc@5 100.00 (100.00)
Test: [70/81] Time 0.017 ( 0.079) Loss 1.1945e+00 (1.2000e+00) Acc@1 50.00 (38.59) Acc@5 100.00 (100.00)
Test: [80/81] Time 0.007 ( 0.077) Loss 1.5687e+00 (1.2169e+00) Acc@1 20.00 (39.63) Acc@5 100.00 (100.00)
* Acc@1 39.630 Acc@5 100.000
```

At the end of the first epoch:

Profiling the model with the following hyperparameter settings: Batch size – 10 Epoch -5 and the Learning Rate set to 0.001.

```
!time python main.py -a alexnet -b 10 --epochs 5 --lr 0.001 Imagenet

Epoch: [0][420/441] Time 0.025 ( 0.069) Data 0.000 ( 0.048) Loss 1.3115e+00 (2.4172e+00) Acc@1 40.00 (24.68) Acc@5 100.00 (96.86)
Epoch: [0][430/441] Time 0.175 ( 0.069) Data 0.167 ( 0.048) Loss 1.3211e+00 (2.3923e+00) Acc@1 30.00 (24.85) Acc@5 100.00 (96.94)
Epoch: [0][440/441] Time 0.131 ( 0.069) Data 0.000 ( 0.048) Loss 1.1985e+00 (2.3684e+00) Acc@1 40.00 (25.18) Acc@5 100.00 (97.00)
Test: [ 0/81] Time 0.855 ( 0.855) Loss 1.5621e+00 (1.5621e+00) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [10/81] Time 0.006 ( 0.125) Loss 1.3568e+00 (1.3213e+00) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [20/81] Time 0.008 ( 0.101) Loss 1.3765e+00 (1.3235e+00) Acc@1 10.00 ( 0.48) Acc@5 100.00 (100.00)
Test: [30/81] Time 0.010 ( 0.089) Loss 1.2447e+00 (1.3082e+00) Acc@1 10.00 ( 6.45) Acc@5 100.00 (100.00)
Test: [40/81] Time 0.009 ( 0.088) Loss 1.2227e+00 (1.3024e+00) Acc@1 40.00 ( 8.78) Acc@5 100.00 (100.00)
Test: [50/81] Time 0.082 ( 0.084) Loss 8.5190e-01 (1.2285e+00) Acc@1 100.00 (25.69) Acc@5 100.00 (100.00)
Test: [60/81] Time 0.086 ( 0.083) Loss 9.7332e-01 (1.1812e+00) Acc@1 80.00 (36.72) Acc@5 100.00 (100.00)
Test: [70/81] Time 0.017 ( 0.079) Loss 1.1945e+00 (1.2000e+00) Acc@1 50.00 (38.59) Acc@5 100.00 (100.00)
Test: [80/81] Time 0.007 ( 0.077) Loss 1.5687e+00 (1.2169e+00) Acc@1 20.00 (39.63) Acc@5 100.00 (100.00)
* Acc@1 39.630 Acc@5 100.000
Epoch: [1][ 0/441] Time 0.620 ( 0.620) Data 0.603 ( 0.603) Loss 1.3165e+00 (1.3165e+00) Acc@1 30.00 (30.00) Acc@5 100.00 (100.00)
Epoch: [1][10/441] Time 0.031 ( 0.112) Data 0.000 ( 0.067) Loss 1.3113e+00 (1.2857e+00) Acc@1 50.00 (39.09) Acc@5 100.00 (100.00)
Epoch: [1][20/441] Time 0.070 ( 0.093) Data 0.062 ( 0.064) Loss 1.2147e+00 (1.2705e+00) Acc@1 30.00 (38.57) Acc@5 100.00 (100.00)
Epoch: [1][30/441] Time 0.119 ( 0.086) Data 0.109 ( 0.060) Loss 1.0326e+00 (1.2791e+00) Acc@1 60.00 (40.32) Acc@5 100.00 (100.00)
```



The time taken to process has significantly increased as compared to the previous two runs as could be expected from the increase in epoch and smaller learning rate. With the decreases in learning rate the step size decreases, and the processing time increases. One can also note that the accuracy has significantly improved for this run in both environments.

Profiling the model with the following hyperparameter settings: Batch size – 10 Epoch -5 and the Learning Rate set to 0.001.

```
!time nvprof python main.py -a alexnet -b 10 --epochs 5 --lr 0.001 Imagenet > log3.txt

==1105== NvPROF is profiling process 1105, command: python3 main.py -a alexnet -b 10 --epochs 5 --lr 0.001 Imagenet
==1105== Profiling application: python3 main.py -a alexnet -b 10 --epochs 5 --lr 0.001 Imagenet
==1105== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities:
41.84%    26.1684s    146288    178.88us    1.0240us    2.0494ms    _ZN2at6native29vectorized_elementwise_kernelIi4EZZNS0_15add_kernel_cudaERNS
9.21%    5.76144s    22039    261.42us    76.511us    932.94us    volta_sgemv_128x64_nt
7.23%    4.52307s    45704    98.964us    1.0240us    1.4574ms    _ZN2at6native29vectorized_elementwise_kernelIi4EZZNS0_23gpu_kernel_with_scala
5.98%    3.73947s    7830    477.58us    104.35us    1.2111ms    volta_sgemv_128x32_sliced1x4_tn
5.58%    3.49381s    12640    276.35us    166.30us    724.28us    volta_sgemv_128x64_nn
3.94%    2.46166s    46289    53.180us    832ns    652.95us    _ZN2at6native29vectorized_elementwise_kernelIi4EZZNS0_16fill_kernel_cudaERN
3.69%    2.30713s    6600    349.57us    80.671us    1.1199ms    volta_sgemv_128x32_sliced1x4_nn
2.44%    1.52304s    5240    290.66us    1.0560us    32.246ms    [CUDA memcopy HtoD]
1.92%    1.20321s    2209    544.68us    292.96us    1.1815ms    void cudnn::detail::wgrad_alg0_engine<float, int=512, int=6, int=5, int=3, in
1.72%    1.07697s    14427    74.649us    58.111us    105.38us    void cudnn::winograd_nonfused::winogradForwardFilter4x4<float, float>(cudnn::
1.52%    0.97502ms    3875    245.68us    1.2160us    99.093ms    [CUDA memcopy DtoH]
1.48%    0.92753ms    14427    64.291us    35.264us    104.93us    void cudnn::winograd_nonfused::winogradForwardData4x4<float, float>(cudnn::wi
1.45%    0.90768ms    14427    62.915us    39.615us    92.383us    void cudnn::winograd_nonfused::winogradForwardOutput4x4<float, float>(cudnn::
1.41%    0.87999ms    33705    26.108us    1.7600us    101.76us    _ZN2at6native29vectorized_elementwise_kernelIi4EZZNS0_21threshold_kernel_impl
1.32%    0.82271ms    2612    314.98us    148.99us    594.33us    volta_cudnn_128x64_relu_xregs_large_nn_v1
1.28%    0.80210ms    6615    121.25us    26.975us    341.85us    void at::native::GLOBAL_N_63_tmpxft_00002580_00000000_10_DilatedMaxPool2d_
1.08%    0.67549ms    6621    102.02us    73.247us    238.65us    void cudnn::winograd_nonfused::winogradForwardOutput4x4<float, float>(cudnn::wi
0.91%    0.57058ms    7026    81.209us    42.408us    228.45us    void cudnn::winograd_nonfused::winogradForwardData9x9_5x5<float, float>(cudnn
0.79%    0.49266ms    4819    102.23us    46.528us    190.56us    void cudnn::winograd_nonfused::winogradForwardOutput9x9_5x5<float, float>(cud
0.73%    0.45732ms    13050    35.043us    11.616us    87.646us    _ZN2at6native27unrolled_elementwise_kernelIiZZNS0_15add_kernel_cudaERNS_14Ten
0.60%    0.37796ms    16245    23.265us    2.1750us    106.85us    _ZN2at6native13reduce_kernelIi5I2EliENS0_8ReduceOpIFNS0_14func_wrapper_tifZ
0.53%    0.33098ms    7830    42.270us    6.7200us    83.871us    void at::native::GLOBAL_N_63_tmpxft_00002580_00000000_10_DilatedMaxPool2d_
0.49%    0.30439ms    6621    45.973us    18.399us    66.976us    void cudnn::winograd_nonfused::winogradForwardDelta4x4<float, float>(cudnn::wi
0.48%    0.30044ms    4819    62.345us    59.487us    73.279us    void cudnn::winograd_nonfused::winogradForwardFilter9x9_5x5<float, float>(cud
0.46%    0.28824ms    2207    130.60us    109.95us    200.73us    void cudnn::winograd_nonfused::winogradForwardDelta9x9_5x5<float, float>(cudnn
0.40%    0.25033ms    6621    37.807us    9.5040us    62.176us    void cudnn::winograd_nonfused::winogradForwardData4x4<float, float>(cudnn::wi
0.24%    0.15267ms    2207    69.177us    63.487us    98.878us    void cudnn::winograd_nonfused::winogradForwardOutput9x9_5x5<float, float>(cudnn
0.20%    0.12601ms    2610    48.279us    24.063us    91.742us    void at::native::GLOBAL_N_69_tmpxft_00001db3_00000000_10_AdaptiveAveragePo
0.16%    0.09563ms    2205    45.194us    25.056us    93.326us    void at::native::GLOBAL_N_69_tmpxft_00001db3_00000000_10_AdaptiveAveragePo
0.08%    0.05249ms    6615    7.9360us    6.3040us    14.271us    _ZN2at6native13reduce_kernelIi5I2EliENS0_8ReduceOpIFNS0_14func_wrapper_tifZ
0.08%    0.05032ms    7830    6.4260us    4.8630us    11.232us    _ZN2at6native27unrolled_elementwise_kernelIiZZNS0_21copy_device_to_deviceERNS
```

```
real    4m22.103s
user    5m45.346s
sys     1m6.378s
```

**Running on BareMetal** : We train the model with the same hyperparameter settings: Batch size – 10 Epoch -5 and the Learning Rate has been set to 0.001. It takes approximately 441 iterations in each epoch and processes images in batches of 10 and passes over the entire dataset 5 times.

```
(py3.6.3) [ai3594@cpu-40 ~]$ time nvprof python /home/ai3594/examples/imagenet/main.py -a alexnet -b 10 --epochs 5 --lr 0.001 /home/ai3594/
--99467-- NvPROF is profiling process 99467, command: python /home/ai3594/examples/imagenet/main.py -a alexnet -b 10 --epochs 5 --lr 0.001 /home/ai3594/
> creating model 'alexnet'
Epoch: [0] 0/441      Time: 0.010 ( 0.010)      Data: 0.629 ( 0.629)      Loss: 6.9115e+00 (6.9115e+00)      Acc@1: 0.00 ( 0.00)      Acc@5: 0.00 ( 0.00)
Epoch: [0] 10/441     Time: 0.031 ( 0.113)      Data: 0.013 ( 0.072)      Loss: 6.9528e+00 (6.9528e+00)      Acc@1: 0.00 ( 0.00)      Acc@5: 0.00 ( 0.00)
Epoch: [0] 20/441     Time: 0.007 ( 0.094)      Data: 0.078 ( 0.064)      Loss: 6.8539e+00 (6.8539e+00)      Acc@1: 50.00 ( 9.05)      Acc@5: 100.00 ( 23.33)
Epoch: [0] 30/441     Time: 0.038 ( 0.103)      Data: 0.000 ( 0.072)      Loss: 6.8179e+00 (6.8179e+00)      Acc@1: 10.00 ( 14.52)      Acc@5: 100.00 ( 46.06)
Epoch: [0] 40/441     Time: 0.041 ( 0.093)      Data: 0.001 ( 0.063)      Loss: 6.7372e+00 (6.7372e+00)      Acc@1: 30.00 ( 17.50)      Acc@5: 100.00 ( 40.73)
Epoch: [0] 50/441     Time: 0.046 ( 0.091)      Data: 0.000 ( 0.061)      Loss: 6.5462e+00 (6.5462e+00)      Acc@1: 40.00 ( 20.00)      Acc@5: 100.00 ( 60.43)
Epoch: [0] 60/441     Time: 0.050 ( 0.089)      Data: 0.000 ( 0.059)      Loss: 6.4607e+00 (6.4607e+00)      Acc@1: 10.00 ( 19.53)      Acc@5: 100.00 ( 73.63)
Epoch: [0] 70/441     Time: 0.039 ( 0.088)      Data: 0.000 ( 0.058)      Loss: 7.3267e+00 (6.4086e+00)      Acc@1: 30.00 ( 20.00)      Acc@5: 100.00 ( 77.32)
Epoch: [0] 80/441     Time: 0.037 ( 0.085)      Data: 0.010 ( 0.055)      Loss: 6.5856e+00 (6.4160e+00)      Acc@1: 30.00 ( 20.74)      Acc@5: 100.00 ( 80.12)
Epoch: [0] 90/441     Time: 0.039 ( 0.085)      Data: 0.000 ( 0.055)      Loss: 6.5496e+00 (6.4340e+00)      Acc@1: 10.00 ( 20.22)      Acc@5: 100.00 ( 82.31)
Epoch: [0] 100/441     Time: 0.040 ( 0.084)      Data: 0.004 ( 0.053)      Loss: 6.4667e+00 (6.4416e+00)      Acc@1: 20.00 ( 20.20)      Acc@5: 100.00 ( 84.06)
Epoch: [0] 110/441     Time: 0.042 ( 0.083)      Data: 0.000 ( 0.053)      Loss: 6.3855e+00 (6.4411e+00)      Acc@1: 40.00 ( 20.99)      Acc@5: 100.00 ( 85.50)
Epoch: [0] 120/441     Time: 0.043 ( 0.083)      Data: 0.001 ( 0.053)      Loss: 5.8615e+00 (6.4203e+00)      Acc@1: 50.00 ( 21.65)      Acc@5: 100.00 ( 86.69)
Epoch: [0] 130/441     Time: 0.039 ( 0.083)      Data: 0.000 ( 0.053)      Loss: 4.3539e+00 (6.2436e+00)      Acc@1: 20.00 ( 22.53)      Acc@5: 100.00 ( 87.71)
Epoch: [0] 140/441     Time: 0.043 ( 0.082)      Data: 0.000 ( 0.052)      Loss: 2.5066e+00 (6.0589e+00)      Acc@1: 20.00 ( 21.56)      Acc@5: 100.00 ( 88.58)
Epoch: [0] 150/441     Time: 0.040 ( 0.082)      Data: 0.000 ( 0.052)      Loss: 2.1422e+00 (5.7984e+00)      Acc@1: 10.00 ( 21.60)      Acc@5: 100.00 ( 89.34)
Epoch: [0] 160/441     Time: 0.041 ( 0.082)      Data: 0.001 ( 0.051)      Loss: 1.7635e+00 (5.5377e+00)      Acc@1: 20.00 ( 21.00)      Acc@5: 100.00 ( 90.00)
Epoch: [0] 170/441     Time: 0.047 ( 0.081)      Data: 0.000 ( 0.051)      Loss: 1.5924e+00 (5.2897e+00)      Acc@1: 20.00 ( 21.64)      Acc@5: 100.00 ( 91.99)
Epoch: [0] 180/441     Time: 0.040 ( 0.082)      Data: 0.000 ( 0.051)      Loss: 1.3950e+00 (5.1036e+00)      Acc@1: 50.00 ( 22.38)      Acc@5: 100.00 ( 91.10)
Epoch: [0] 190/441     Time: 0.036 ( 0.081)      Data: 0.005 ( 0.051)      Loss: 1.3510e+00 (4.9151e+00)      Acc@1: 40.00 ( 22.53)      Acc@5: 100.00 ( 91.57)
Epoch: [0] 200/441     Time: 0.041 ( 0.082)      Data: 0.002 ( 0.052)      Loss: 1.1548e+00 (4.7447e+00)      Acc@1: 10.00 ( 22.44)      Acc@5: 100.00 ( 91.99)
Epoch: [0] 210/441     Time: 0.163 ( 0.081)      Data: 0.154 ( 0.052)      Loss: 1.8661e+00 (4.5852e+00)      Acc@1: 10.00 ( 22.84)      Acc@5: 100.00 ( 92.37)
Epoch: [0] 220/441     Time: 0.041 ( 0.081)      Data: 0.004 ( 0.051)      Loss: 1.2739e+00 (4.4447e+00)      Acc@1: 50.00 ( 22.74)      Acc@5: 100.00 ( 92.71)
Epoch: [0] 230/441     Time: 0.164 ( 0.081)      Data: 0.155 ( 0.051)      Loss: 1.3979e+00 (4.3217e+00)      Acc@1: 30.00 ( 22.90)      Acc@5: 100.00 ( 93.03)
Epoch: [0] 240/441     Time: 0.040 ( 0.080)      Data: 0.000 ( 0.050)      Loss: 1.2848e+00 (4.2000e+00)      Acc@1: 20.00 ( 22.88)      Acc@5: 100.00 ( 93.22)
Epoch: [0] 250/441     Time: 0.132 ( 0.080)      Data: 0.124 ( 0.050)      Loss: 1.3864e+00 (4.0972e+00)      Acc@1: 30.00 ( 22.95)      Acc@5: 100.00 ( 93.59)
Epoch: [0] 260/441     Time: 0.039 ( 0.080)      Data: 0.005 ( 0.051)      Loss: 1.3677e+00 (3.9947e+00)      Acc@1: 30.00 ( 23.07)      Acc@5: 100.00 ( 93.83)
Epoch: [0] 270/441     Time: 0.145 ( 0.080)      Data: 0.136 ( 0.051)      Loss: 1.4539e+00 (3.9017e+00)      Acc@1: 30.00 ( 22.92)      Acc@5: 100.00 ( 94.09)
Epoch: [0] 280/441     Time: 0.046 ( 0.080)      Data: 0.004 ( 0.050)      Loss: 1.4913e+00 (3.8151e+00)      Acc@1: 30.00 ( 22.99)      Acc@5: 100.00 ( 94.27)
Epoch: [0] 290/441     Time: 0.142 ( 0.080)      Data: 0.133 ( 0.050)      Loss: 1.4231e+00 (3.7146e+00)      Acc@1: 10.00 ( 22.74)      Acc@5: 100.00 ( 94.47)
Epoch: [0] 300/441     Time: 0.053 ( 0.079)      Data: 0.000 ( 0.049)      Loss: 1.4098e+00 (3.6602e+00)      Acc@1: 10.00 ( 22.99)      Acc@5: 100.00 ( 94.65)
Epoch: [0] 310/441     Time: 0.049 ( 0.079)      Data: 0.000 ( 0.050)      Loss: 1.3719e+00 (3.5950e+00)      Acc@1: 10.00 ( 22.99)      Acc@5: 100.00 ( 94.82)
Epoch: [0] 320/441     Time: 0.040 ( 0.079)      Data: 0.005 ( 0.050)      Loss: 1.3422e+00 (3.5237e+00)      Acc@1: 40.00 ( 22.99)      Acc@5: 100.00 ( 94.99)
Epoch: [0] 330/441     Time: 0.174 ( 0.079)      Data: 0.165 ( 0.050)      Loss: 1.2946e+00 (3.4587e+00)      Acc@1: 30.00 ( 22.90)      Acc@5: 100.00 ( 95.14)
Epoch: [0] 340/441     Time: 0.036 ( 0.080)      Data: 0.000 ( 0.050)      Loss: 1.3822e+00 (3.4046e+00)      Acc@1: 40.00 ( 23.28)      Acc@5: 100.00 ( 95.28)
Epoch: [0] 350/441     Time: 0.050 ( 0.079)      Data: 0.010 ( 0.050)      Loss: 1.7423e+00 (3.3496e+00)      Acc@1: 40.00 ( 23.45)      Acc@5: 100.00 ( 95.41)
Epoch: [0] 360/441     Time: 0.037 ( 0.080)      Data: 0.002 ( 0.050)      Loss: 1.4713e+00 (3.2961e+00)      Acc@1: 20.00 ( 23.74)      Acc@5: 100.00 ( 95.54)
Epoch: [0] 370/441     Time: 0.177 ( 0.080)      Data: 0.168 ( 0.051)      Loss: 1.7681e+00 (3.2490e+00)      Acc@1: 10.00 ( 23.75)      Acc@5: 100.00 ( 95.66)
Epoch: [0] 380/441     Time: 0.037 ( 0.080)      Data: 0.000 ( 0.050)      Loss: 1.3679e+00 (3.1926e+00)      Acc@1: 30.00 ( 23.57)      Acc@5: 100.00 ( 95.77)
Epoch: [0] 390/441     Time: 0.156 ( 0.080)      Data: 0.147 ( 0.051)      Loss: 1.3274e+00 (3.1570e+00)      Acc@1: 50.00 ( 23.68)      Acc@5: 100.00 ( 95.80)
Epoch: [0] 400/441     Time: 0.036 ( 0.080)      Data: 0.002 ( 0.050)      Loss: 1.3037e+00 (3.1142e+00)      Acc@1: 30.00 ( 23.69)      Acc@5: 100.00 ( 95.89)
Epoch: [0] 410/441     Time: 0.092 ( 0.080)      Data: 0.072 ( 0.051)      Loss: 1.3265e+00 (3.0730e+00)      Acc@1: 40.00 ( 23.70)      Acc@5: 100.00 ( 96.00)
Epoch: [0] 420/441     Time: 0.040 ( 0.079)      Data: 0.002 ( 0.050)      Loss: 1.5491e+00 (3.0344e+00)      Acc@1: 10.00 ( 23.54)      Acc@5: 100.00 ( 96.10)
Epoch: [0] 430/441     Time: 0.072 ( 0.079)      Data: 0.023 ( 0.051)      Loss: 1.4967e+00 (2.9969e+00)      Acc@1: 40.00 ( 23.74)      Acc@5: 100.00 ( 96.20)
Epoch: [0] 440/441     Time: 0.114 ( 0.079)      Data: 0.000 ( 0.050)      Loss: 1.3767e+00 (2.9630e+00)      Acc@1: 0.00 ( 23.77)      Acc@5: 100.00 ( 96.35)
Test: [0/80]      Time: 0.072 ( 0.072)      Loss: 1.4310e+00 (1.4310e+00)      Acc@1: 0.00 ( 0.00)      Acc@5: 100.00 (100.00)
Test: [10/80]     Time: 0.037 ( 0.126)      Loss: 1.4492e+00 (1.4442e+00)      Acc@1: 0.00 ( 0.00)      Acc@5: 100.00 (100.00)
Test: [20/80]     Time: 0.004 ( 0.111)      Loss: 1.2932e+00 (1.4422e+00)      Acc@1: 0.00 ( 0.00)      Acc@5: 100.00 (100.00)
Test: [30/80]     Time: 0.021 ( 0.097)      Loss: 1.4004e+00 (1.4275e+00)      Acc@1: 0.00 ( 0.00)      Acc@5: 100.00 (100.00)
```

```

Epoch: [0][400/441] Time 0.036 ( 0.080) Data 0.002 ( 0.050) Loss 1.3017e+00 (3.1142e+00) Acc@1 50.00 ( 23.67) Acc@5 100.00 ( 95.99)
Epoch: [0][410/441] Time 0.082 ( 0.080) Data 0.072 ( 0.051) Loss 1.3265e+00 (3.0730e+00) Acc@1 40.00 ( 23.70) Acc@5 100.00 ( 96.08)
Epoch: [0][420/441] Time 0.040 ( 0.079) Data 0.003 ( 0.050) Loss 1.5491e+00 (3.0344e+00) Acc@1 10.00 ( 23.54) Acc@5 100.00 ( 96.18)
Epoch: [0][430/441] Time 0.032 ( 0.079) Data 0.023 ( 0.051) Loss 1.4967e+00 (2.9969e+00) Acc@1 40.00 ( 23.74) Acc@5 100.00 ( 96.26)
Epoch: [0][440/441] Time 0.114 ( 0.079) Data 0.000 ( 0.050) Loss 1.3767e+00 (2.9630e+00) Acc@1 0.00 ( 23.77) Acc@5 100.00 ( 96.35)
Test: [ 0/80] Time 0.872 ( 0.872) Loss 1.4510e+00 (1.4510e+00) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [10/80] Time 0.017 ( 0.126) Loss 1.4492e+00 (1.4442e+00) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [20/80] Time 0.004 ( 0.111) Loss 1.3928e+00 (1.4422e+00) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [30/80] Time 0.021 ( 0.097) Loss 1.4004e+00 (1.4275e+00) Acc@1 0.00 ( 0.00) Acc@5 100.00 (100.00)
Test: [40/80] Time 0.005 ( 0.097) Loss 1.3207e+00 (1.4180e+00) Acc@1 80.00 ( 2.20) Acc@5 100.00 (100.00)
Test: [50/80] Time 0.017 ( 0.090) Loss 1.3222e+00 (1.3990e+00) Acc@1 70.00 ( 16.67) Acc@5 100.00 (100.00)
Test: [60/80] Time 0.005 ( 0.091) Loss 1.3366e+00 (1.3864e+00) Acc@1 90.00 ( 26.72) Acc@5 100.00 (100.00)
Test: [70/80] Time 0.026 ( 0.087) Loss 1.3391e+00 (1.3818e+00) Acc@1 90.00 ( 35.92) Acc@5 100.00 (100.00)
* Acc@1 41.375 Acc@5 100.000

```

At the end of the first epoch: One can observe the accuracy is 41.375. Once all five epochs' have been processed and the training is complete. The time taken to process has significantly increased as compared to the previous two runs as could be expected from the increase in epoch and smaller learning rate. With the decreases in learning rate the step size decreases, and the processing time increases. One can also note that the accuracy has significantly improved for this run due to hyperparameter tuning to 78%.

```

Epoch: [4][390/441] Time 0.038 ( 0.079) Data 0.000 ( 0.049) Loss 7.8238e-01 (6.8428e-01) Acc@1 70.00 ( 74.76) Acc@5 100.00 (100.00)
Epoch: [4][400/441] Time 0.042 ( 0.079) Data 0.005 ( 0.049) Loss 8.4305e-01 (6.8319e-01) Acc@1 70.00 ( 74.74) Acc@5 100.00 (100.00)
Epoch: [4][410/441] Time 0.037 ( 0.079) Data 0.000 ( 0.049) Loss 5.6928e-01 (6.8412e-01) Acc@1 80.00 ( 74.77) Acc@5 100.00 (100.00)
Epoch: [4][420/441] Time 0.044 ( 0.079) Data 0.003 ( 0.048) Loss 5.7127e-01 (6.7821e-01) Acc@1 80.00 ( 75.08) Acc@5 100.00 (100.00)
Epoch: [4][430/441] Time 0.044 ( 0.079) Data 0.000 ( 0.048) Loss 5.1225e-01 (6.7372e-01) Acc@1 80.00 ( 75.29) Acc@5 100.00 (100.00)
Epoch: [4][440/441] Time 0.025 ( 0.078) Data 0.000 ( 0.048) Loss 1.4669e+00 (6.7384e-01) Acc@1 80.00 ( 75.26) Acc@5 100.00 (100.00)
Test: [ 0/80] Time 0.850 ( 0.850) Loss 2.6965e-02 (2.6965e-02) Acc@1 100.00 (100.00) Acc@5 100.00 (100.00)
Test: [10/80] Time 0.019 ( 0.126) Loss 1.5859e-02 (5.9705e-02) Acc@1 100.00 ( 96.36) Acc@5 100.00 (100.00)
Test: [20/80] Time 0.004 ( 0.111) Loss 5.7692e-02 (7.3932e-02) Acc@1 100.00 ( 96.19) Acc@5 100.00 (100.00)
Test: [30/80] Time 0.020 ( 0.096) Loss 3.7044e-02 (8.2952e-02) Acc@1 100.00 ( 95.81) Acc@5 100.00 (100.00)
Test: [40/80] Time 0.005 ( 0.096) Loss 3.4009e-01 (1.1921e-01) Acc@1 90.00 ( 95.12) Acc@5 100.00 (100.00)
Test: [50/80] Time 0.025 ( 0.090) Loss 1.0150e+00 (2.3893e-01) Acc@1 80.00 ( 91.96) Acc@5 100.00 (100.00)
Test: [60/80] Time 0.005 ( 0.090) Loss 1.0152e+00 (3.2801e-01) Acc@1 50.00 ( 89.02) Acc@5 100.00 (100.00)
Test: [70/80] Time 0.021 ( 0.086) Loss 1.4189e+00 (4.9465e-01) Acc@1 40.00 ( 82.25) Acc@5 100.00 (100.00)
* Acc@1 78.375 Acc@5 100.000
==95467== Profiling application: python /home/as13594/examples/imagenet/main.py -a alexnet -b 10 --epochs 5 --lr 0.001 /home/as13594/
==95467== Profiling result:
Type      Time      Calls      Avg      Min      Max      Name
GPU activities: 44.76% 26.3893s 146278 180.40us 1.0240us 2.0138ms _ZN2at6native29vectorized_elementwise_kernelILi4EZZZNS0_15add_kernel_cudaE
lve_clEvENKulvE2_clEvEulffE NS 6detail5ArrayIPcLi3EEEEviT0 Tl
7.48% 4.40744s 45684 96.476us 1.0560us 1.3687ms _ZN2at6native29vectorized_elementwise_kernelILi4EZZNS0_23gpu_kernel_with_sc
ensorIteratorEENKulvE2_clEvEulffE EEvs4 RKT EulffE NS 6detail5ArrayIPcLi2EEEEviT0 Tl
5.41% 3.19010s 7815 408.20us 104.52us 816.96us sgemm 32x32x32 NT vec
5.18% 3.05597s 14406 212.13us 179.24us 306.41us maxwell_scudnn_winograd 128x128 ldg1 ldg4 tile228n nt
4.61% 2.71661s 6615 410.67us 78.851us 667.96us void at::native::GLOBAL_N_63tmpxft_000019f5_00000000_10_DilatedMaxPool
_dilatedMaxPool2d_compute_75_cppl_i_db999de0::max_pool_backward_nchw<float, float>*(int, float const *, long const *, int, int, int, int, int, int, int, at::native::GLO
0.95% 562.41ms 2607 215.73us 108.52us 248.01us maxwell_scudnn_128x64_relu_large_nn
0.92% 543.91ms 14442 37.661us 27.361us 52.418us void cudnn::winograd::generateWinogradTilesKernel<int=0, float, float>(cud
ams<float, float>)

```

```

real    4m34.781s
user    3m52.378s
sys     0m33.779s

```

When we compare the runtime of this job on the BareMetal and Cloud Environments, we notice that the time elapsed in both cases was relatively same. The relative peak in GPU time activity on was also comparable in both cases. The throughput or the number of instructions processed per second is also relatively same.

**Conclusion:** The choice between using BareMetal and Cloud based solutions depends a lot on the kind of application and data we need to process. In case the application needs scalable, easy to use, secure and variable workloads with a focus towards better resource utilization, cloud environments are preferable. If the focus is on highly optimizable secure environments that offer low network latency and dedicated resource utilization the outlook is better with BareMetal Solutions.